

Byzantine Paxos

(From Paxos to PBFT)

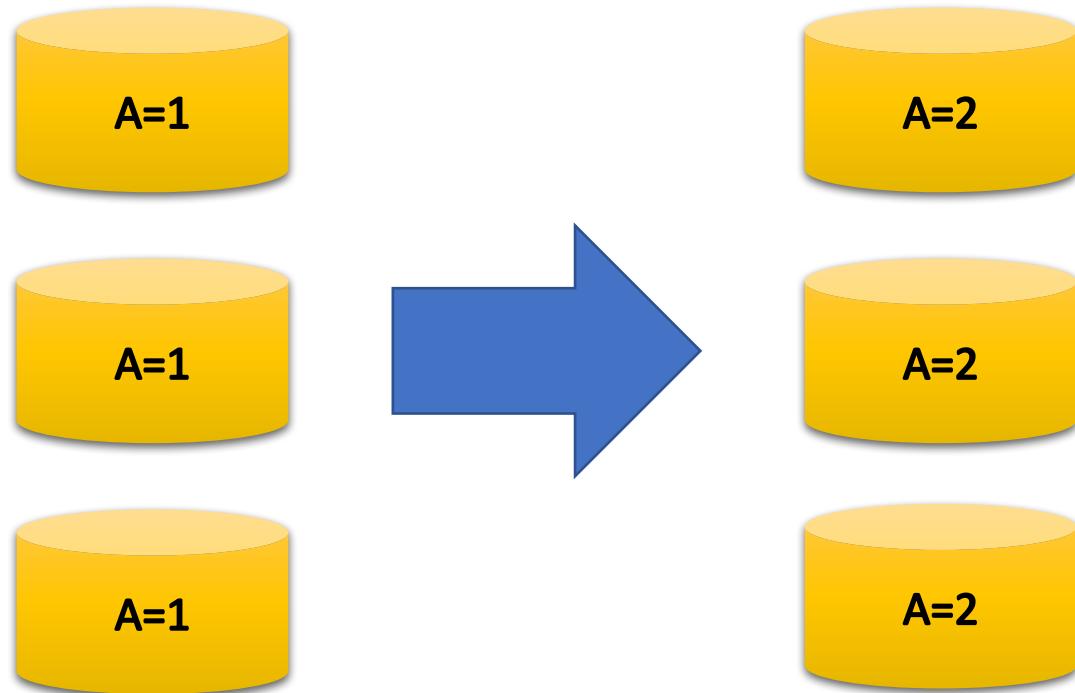
Presenter: Shuai Mu

Structure

- **Background**
 - Replication and consensus
 - Challenges
 - Byzantine failures
- Paxos Review
- ByzPaxos (Non-leader lies)
- ByzPaxos (Leader lies)

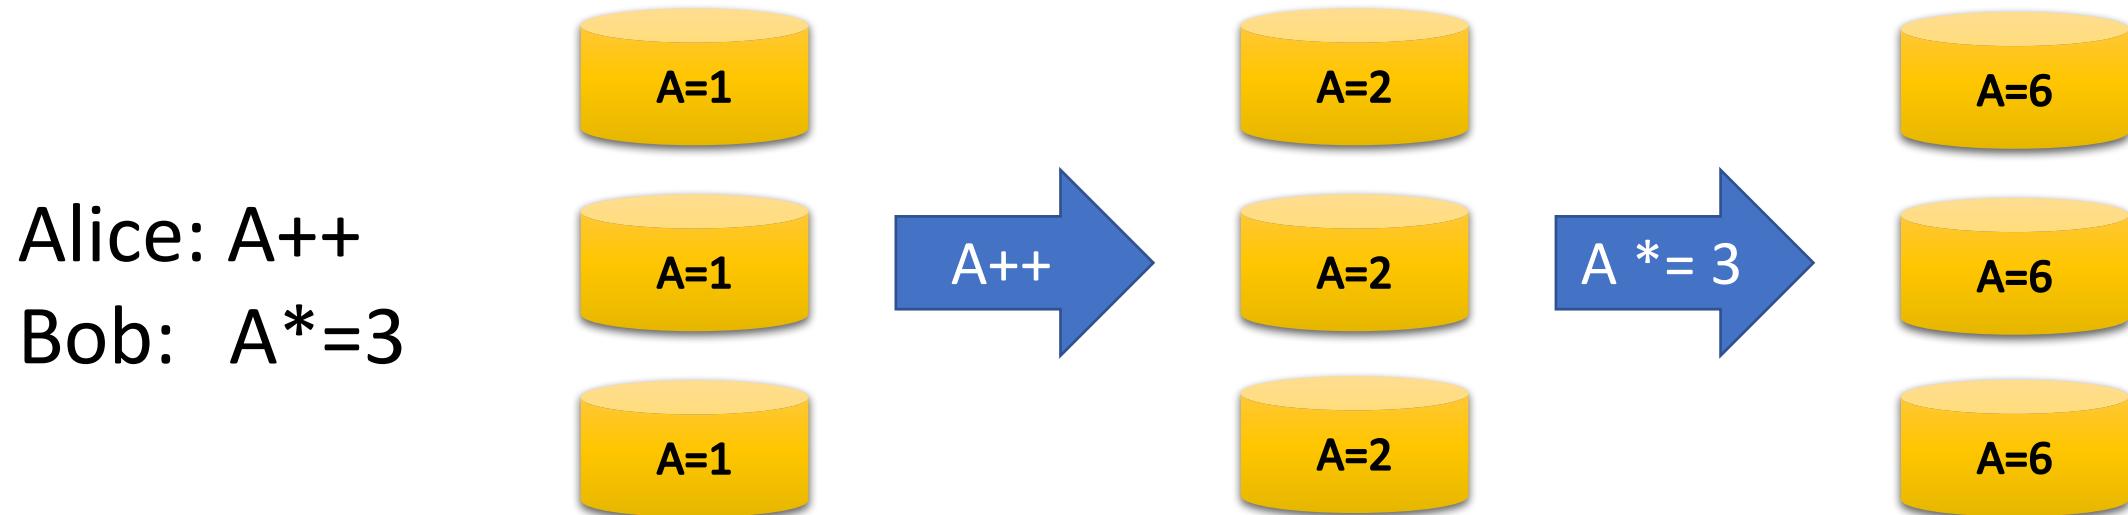
Replication and consensus

- Replication is usual way to achieve reliability
- Each replica must be consistent
- Safety and liveness with a number of failures



Replicated state machine

- Each step must be identical across replicas
- you can make one decision, iff. you can make a series of decisions



Why making a decision is difficult?

- Asynchronous
 - Any network partitions
- Fault tolerance (Crashes)
 - Hopefully tolerate f crashes with $N = 2f + 1$ replicas
- Nobody has a God view (there is no almighty leader)
 - Cannot detect failures
- A Classic Solution: Paxos

Let us makes it even harder!

- Byzantine Failures
 - Beyond crashes, any message is possible
- They actually happen:
 - Malicious attackers
 - Unexpected failures in network, disks, etc.
- What we cannot protect
 - A malicious leader following the protocol issuing a harmful command, e.g. delete a table

Restrict BFT failures

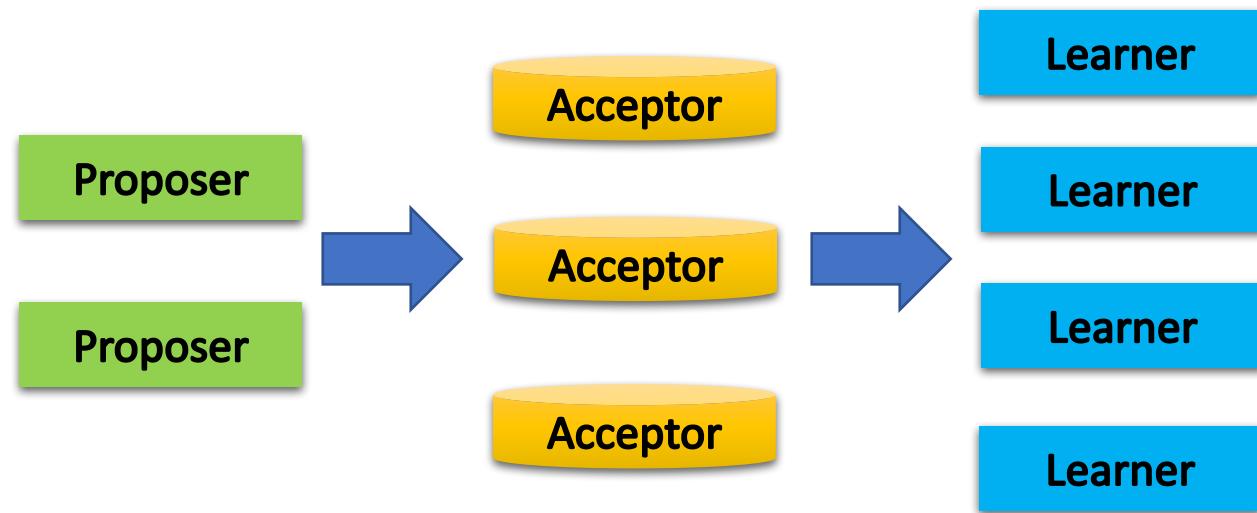
- The problem is too difficult with arbitrary forging.
 - If anyone can pretend to be anyone else, then it looks like anyone is faulty
- Disable forging by signing.
 - Everyone knows public key of everyone else
 - Everyone signs its message with private key

Structure

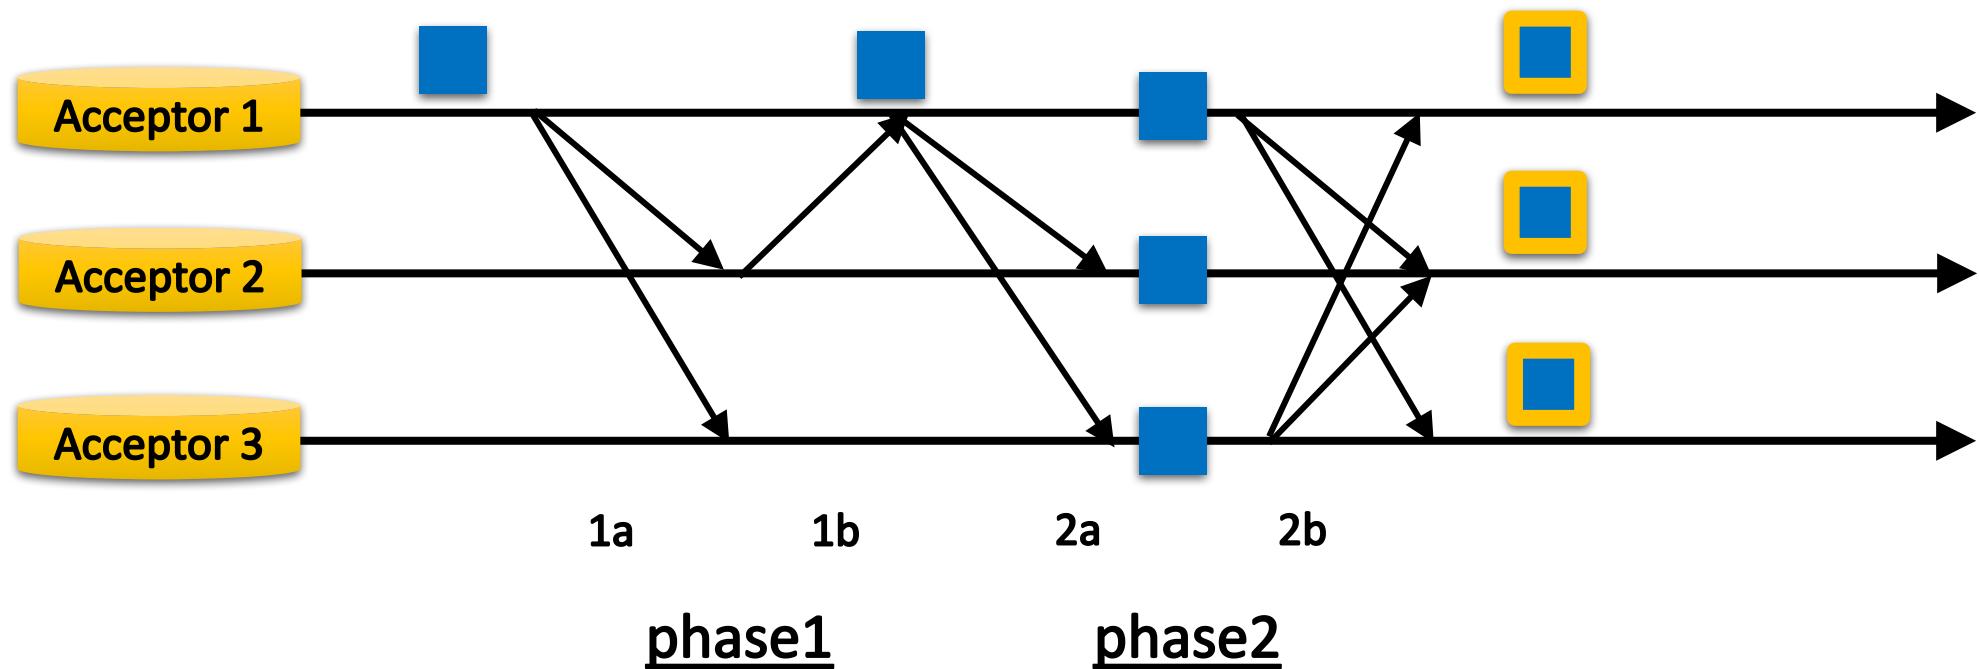
- Background
- **Paxos Review**
 - Overview
 - Protocols
 - Example
- ByzPaxos (Non-leader lies)
- ByzPaxos (Leader lies)

Paxos Review

- Problem
 - Make a decision
- Roles
 - Proposer
 - Acceptor
 - Learner



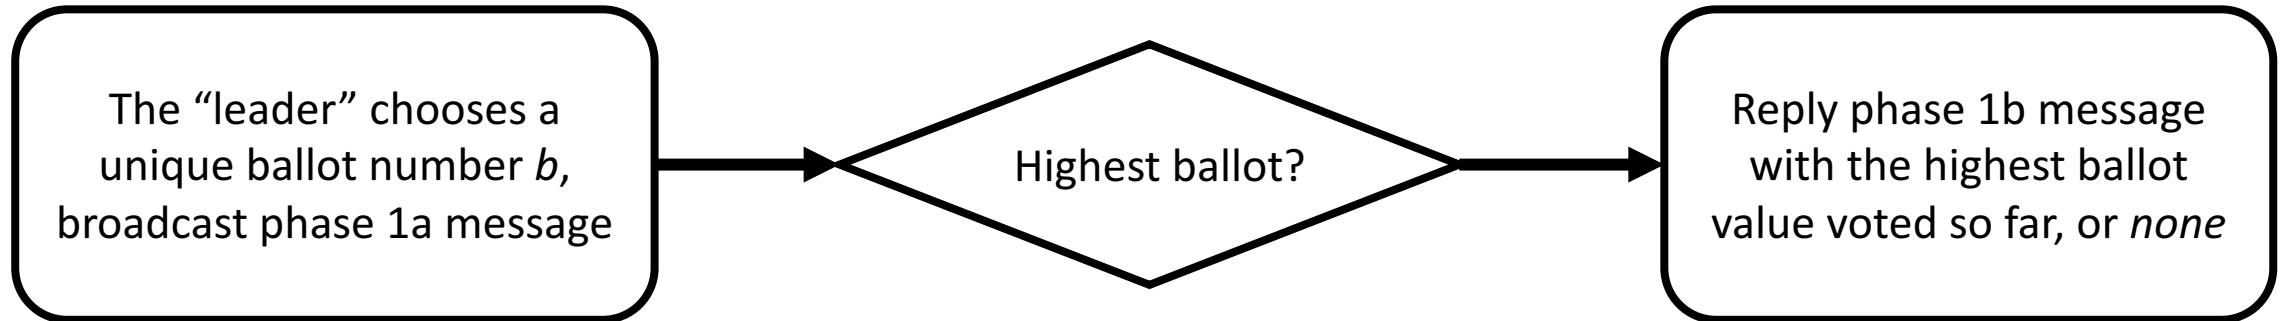
Protocol Overview



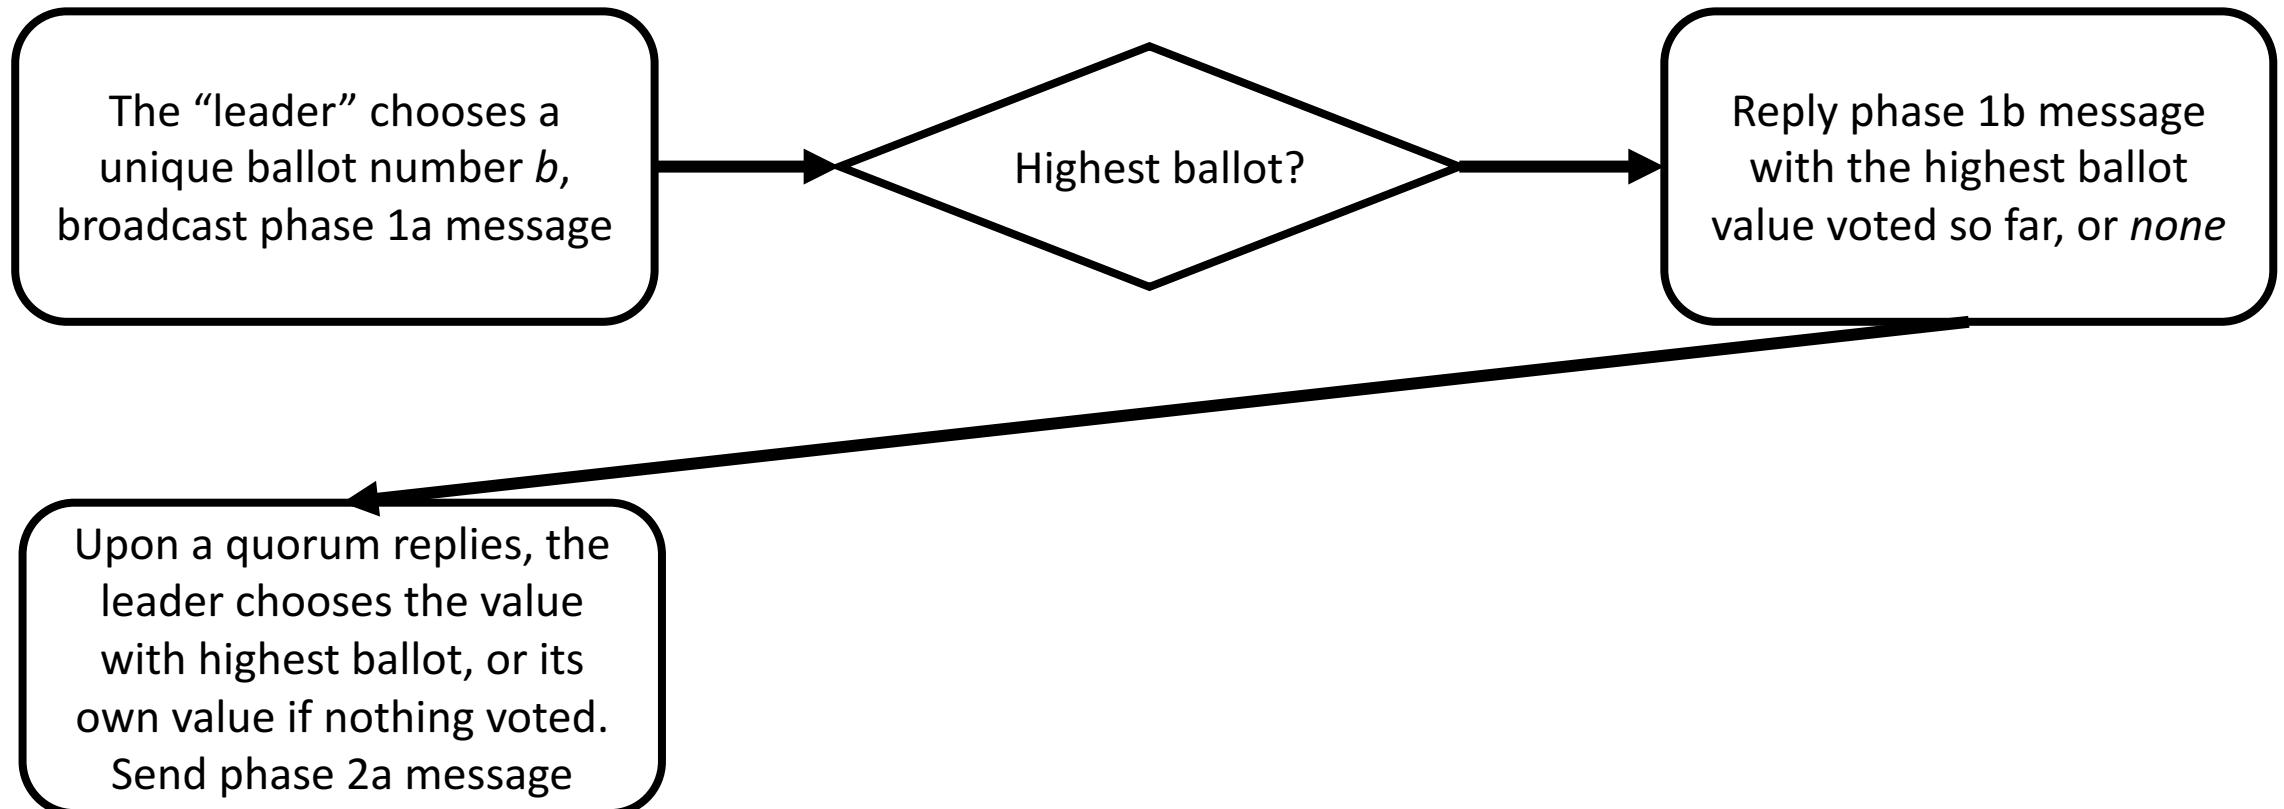
Phase 1a-1b-2a-2b

The “leader” chooses a unique ballot number b , broadcast phase 1a message

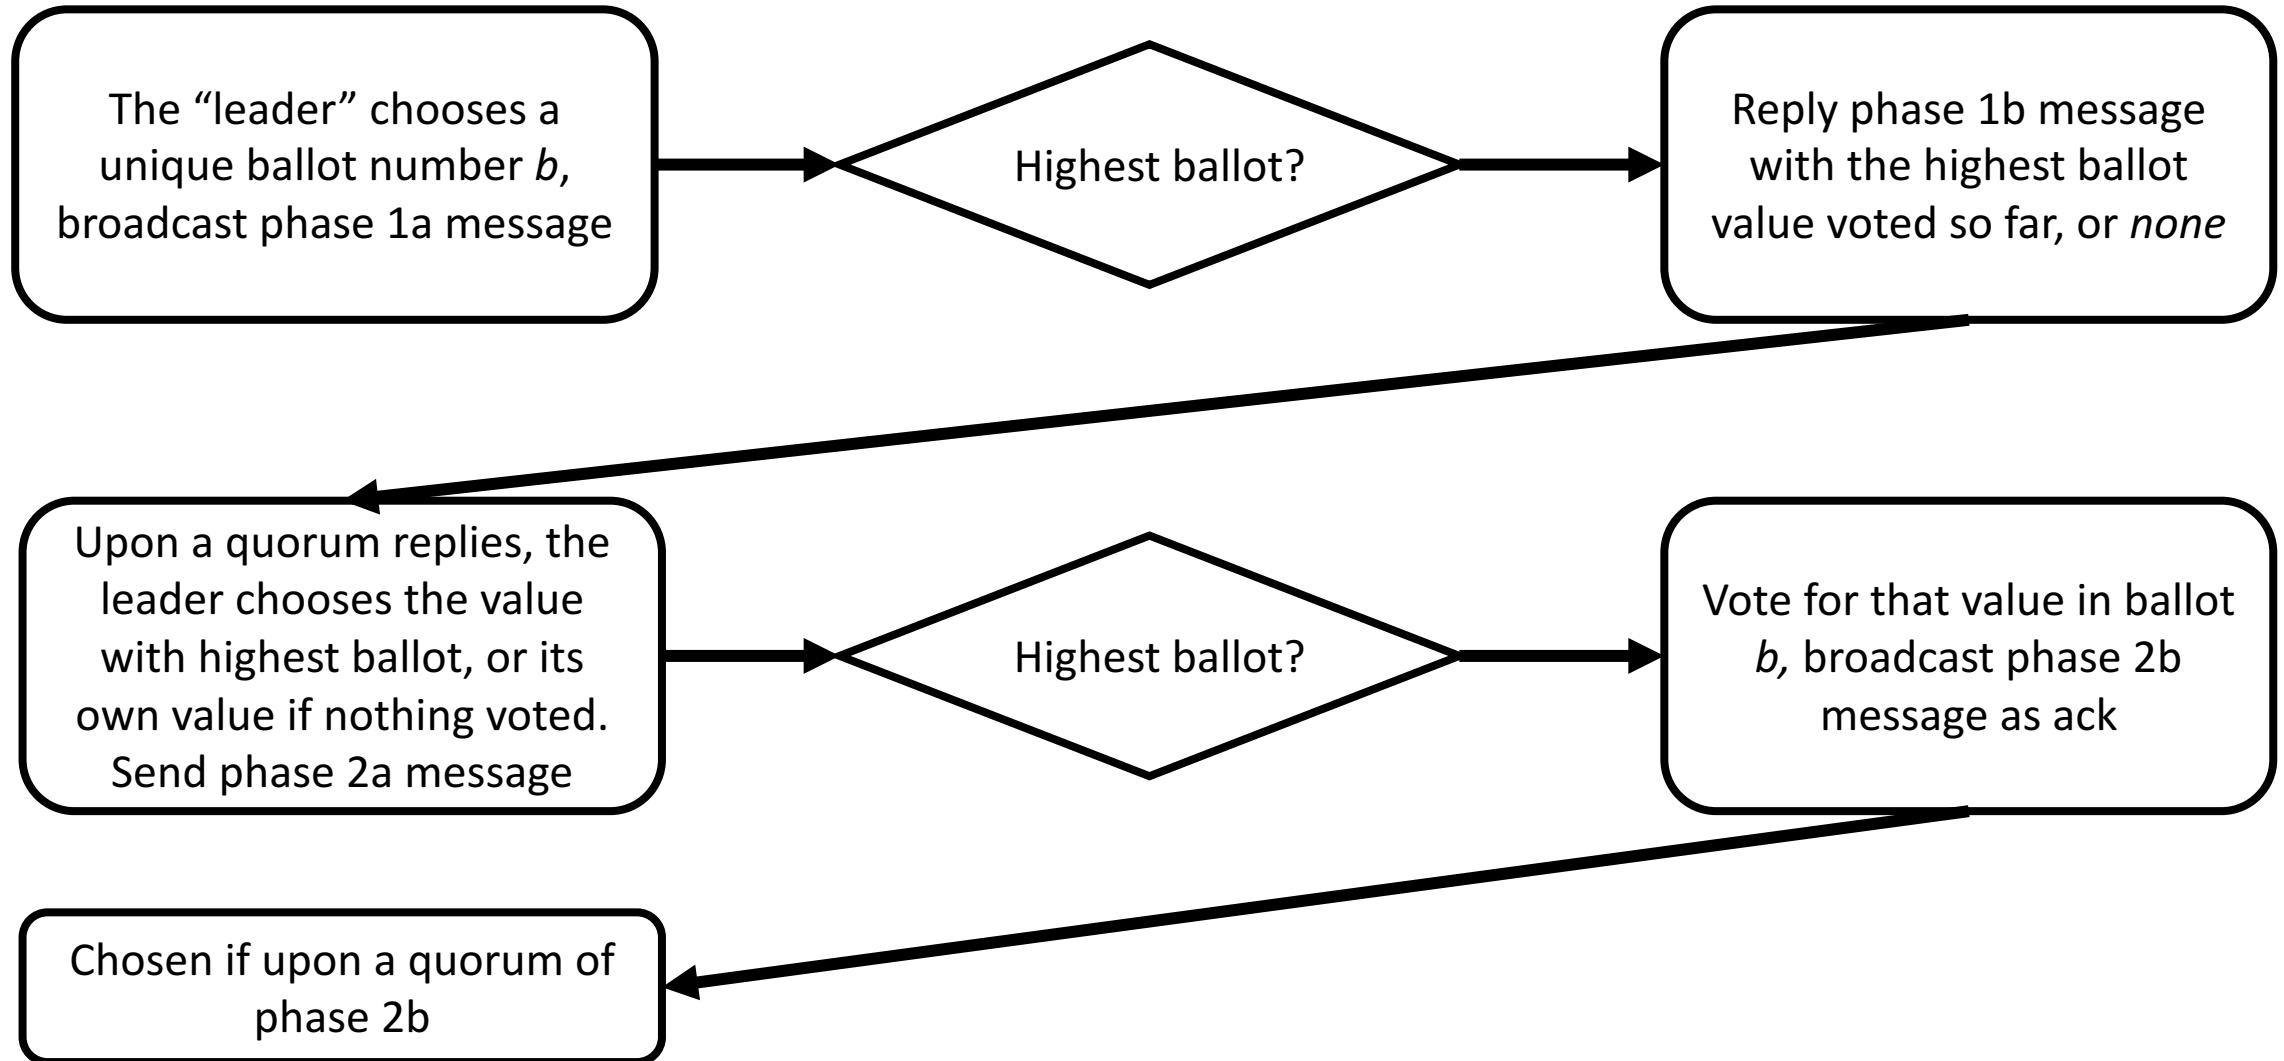
Phase 1a-**1b**-2a-2b



Phase 1a-1b-**2a**-2b

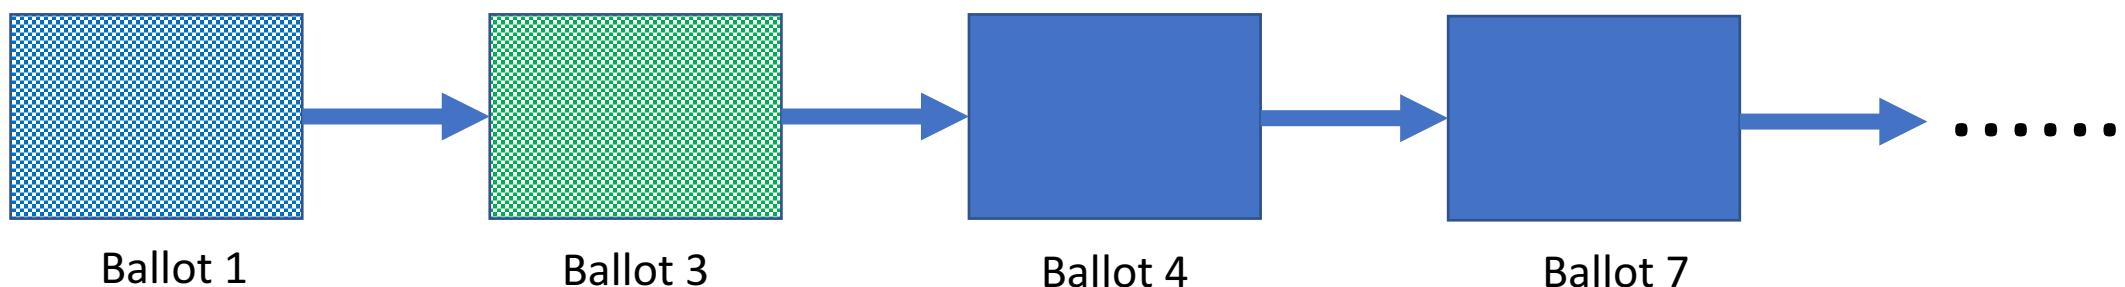


Phase 1a-1b-2a-2b



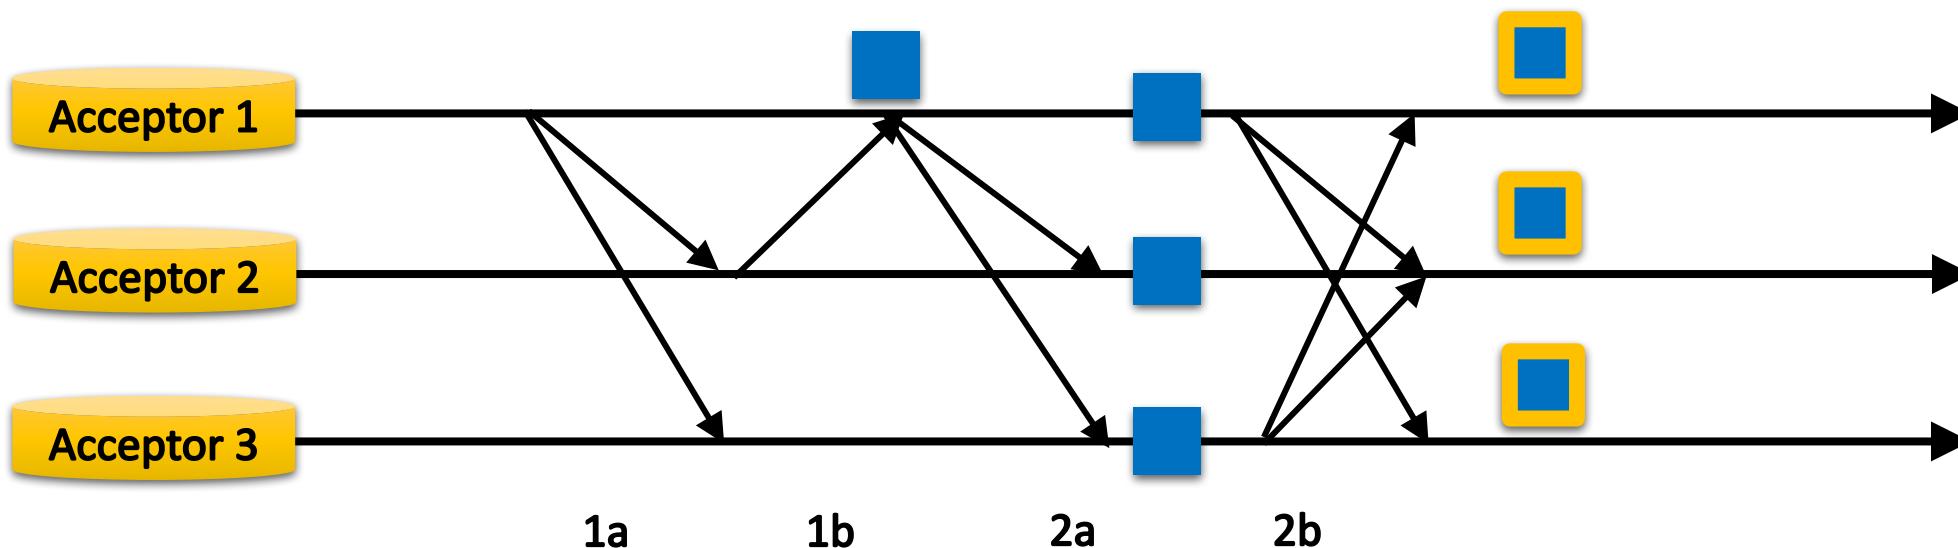
How it works intuitively

- Different acceptors cannot vote for different values in the same ballot.
- An acceptor can vote for different values in different ballots.
- An acceptor can vote for a value iff no other value was(will be) chosen previously

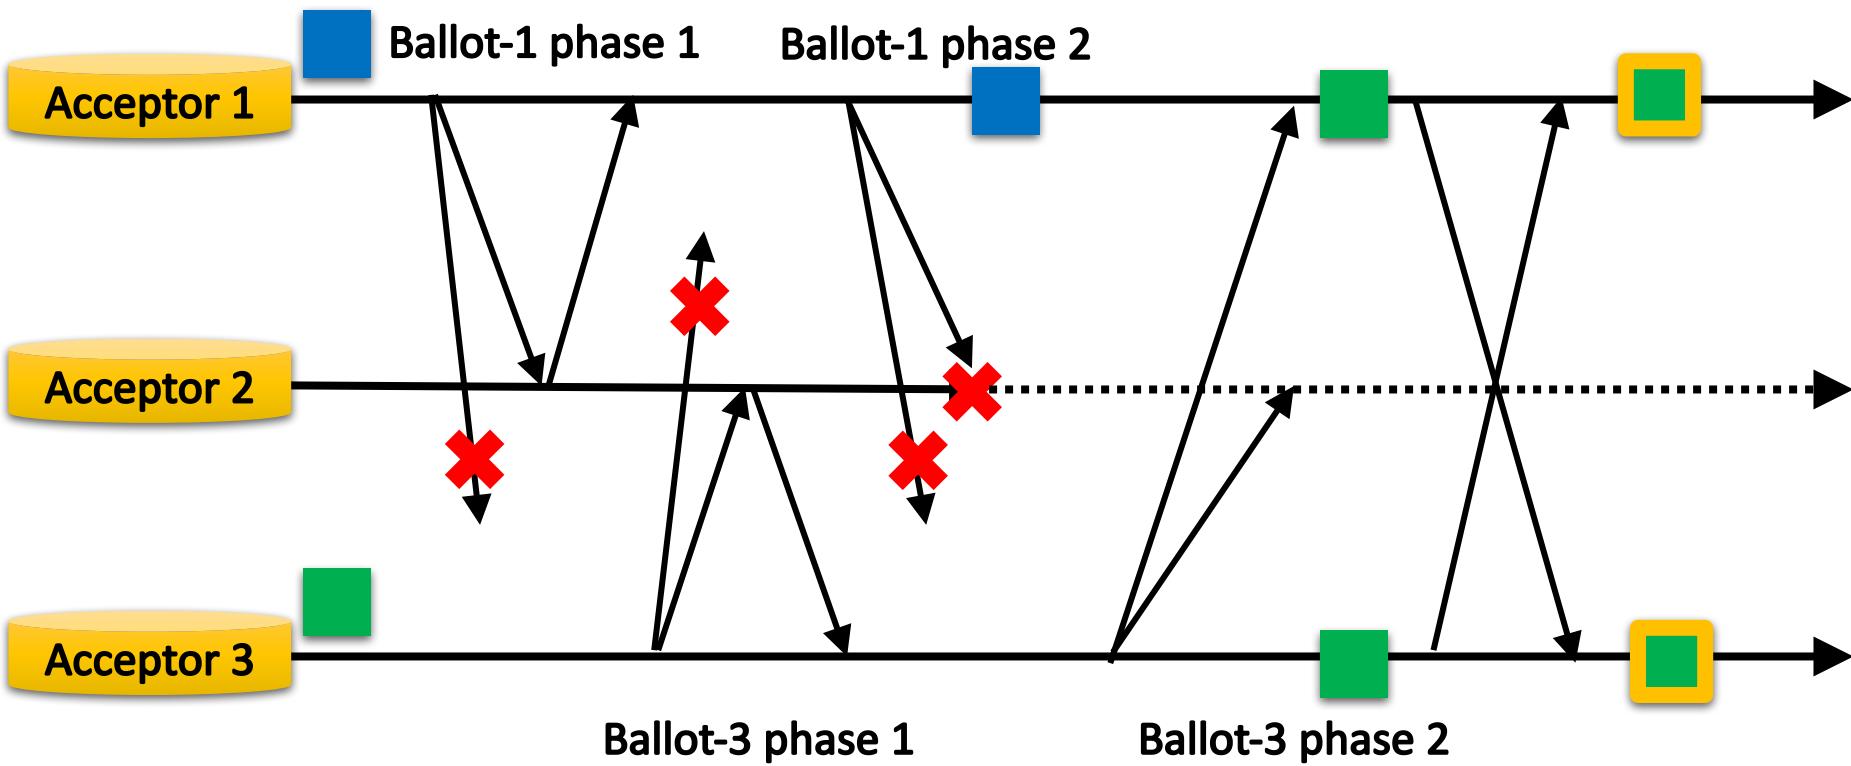


Magic happens between phase 1b and 2a

- If no acceptor in the quorum has voted in a ballot numbered less than b , then all values are safe at b .
- If some acceptor in the quorum has voted, let c be the highest-numbered ballot less than b in which such a vote was cast. The value voted for in ballot c is safe at b .



Example

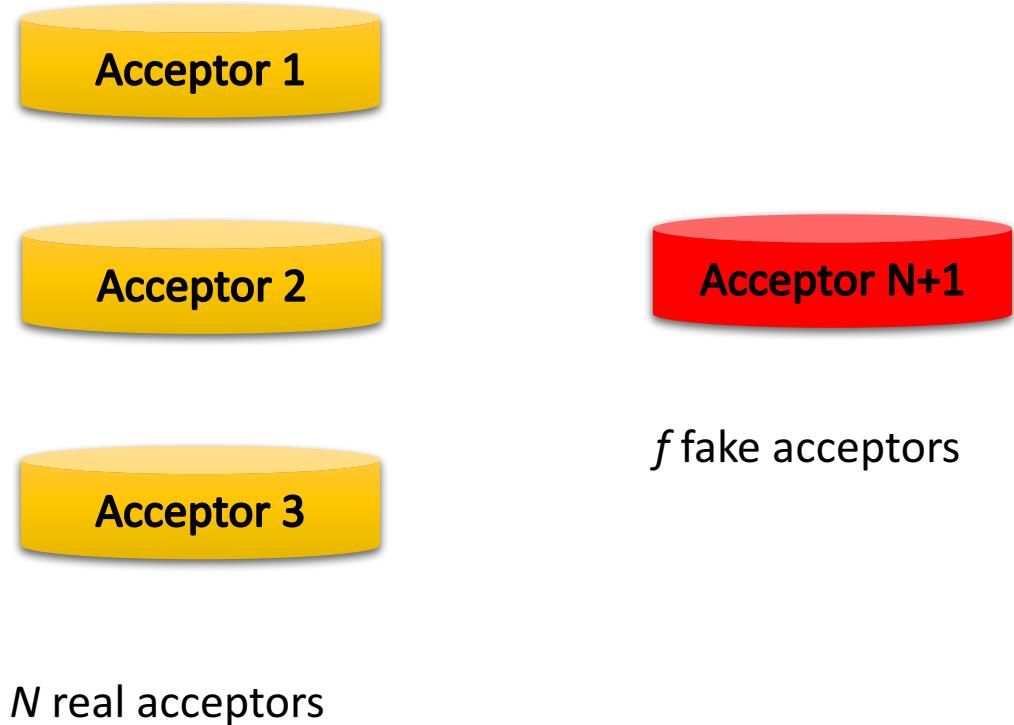


Structure

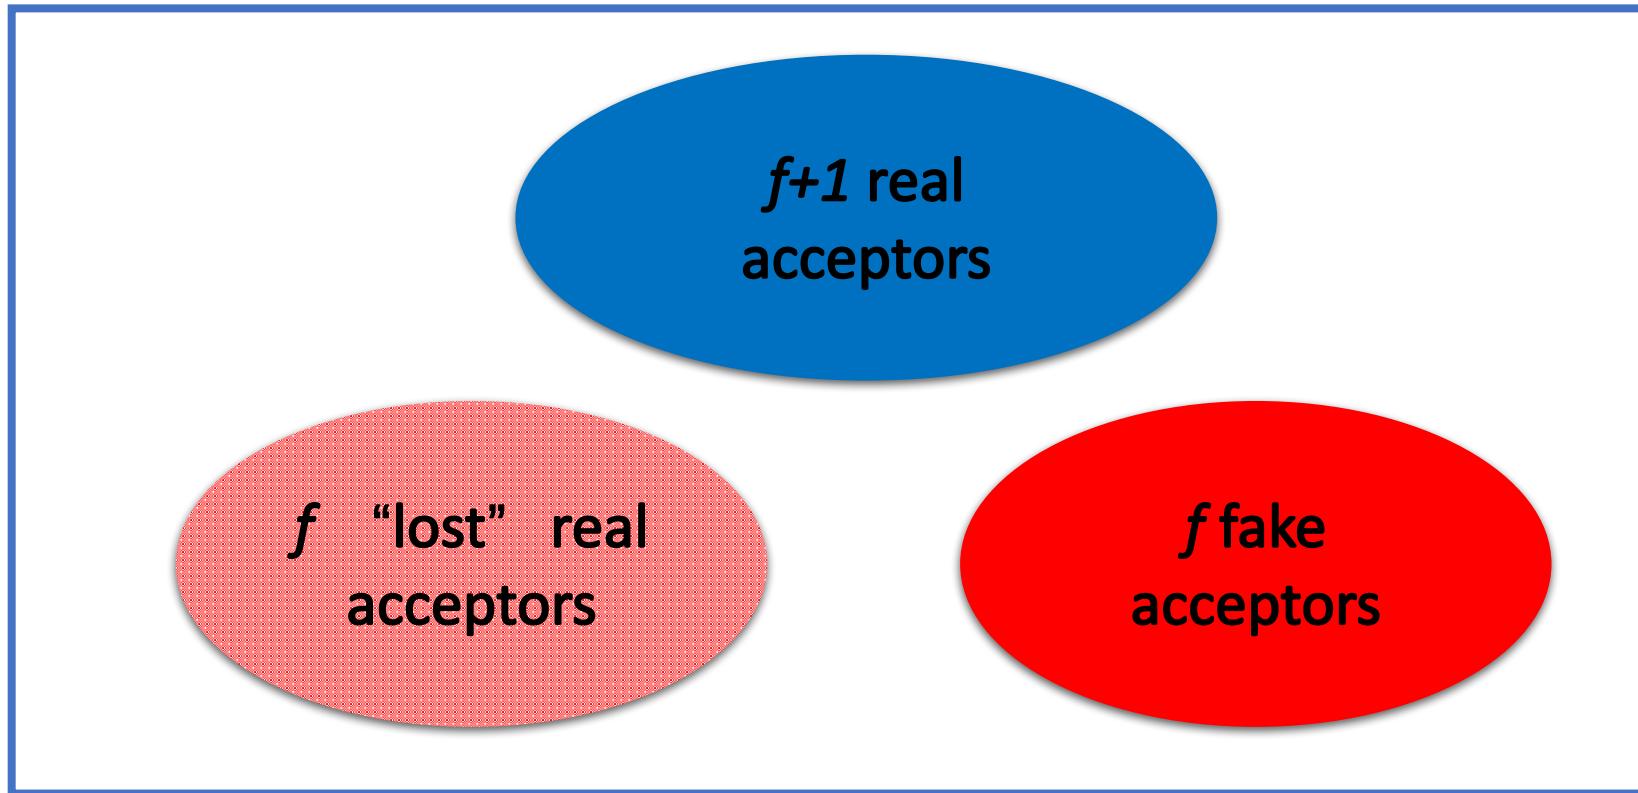
- Background
- Paxos Review
- **ByzPaxos (Non-leader lies)**
 - Goal
 - Solution w/ larger quorum
 - Reduce quorum size
- ByzPaxos (Leader lies)

Let us bring in fake acceptors

- $N+f$ byz-acceptors
 - N real acceptors
 - f fake acceptors,
 - Non-distinguishable
- Consensus?
 - If so, what size is N ?
 - N , in the best case, is $2f+1$, which means $N'=N+f=3f+1$

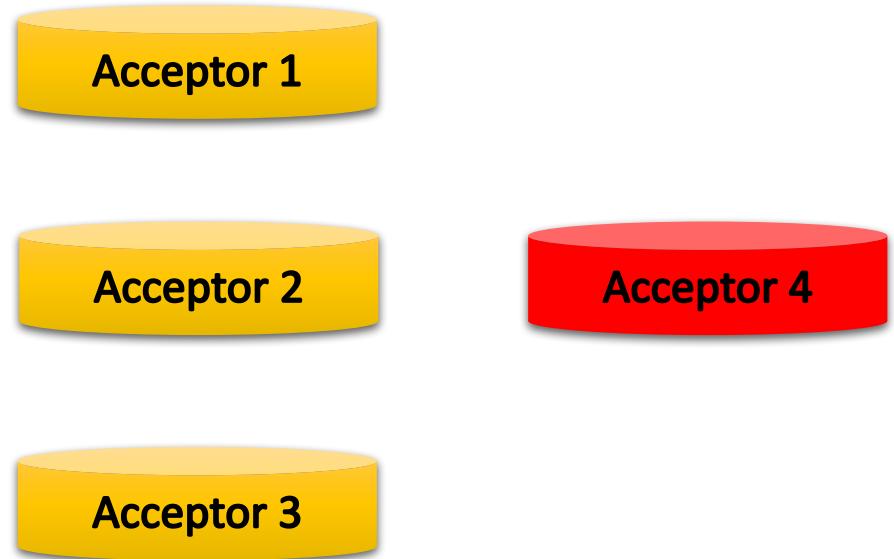


Why at least $3f+1$ acceptors?



Naïve solution: bring extra f into Paxos

- New byz-quorum $q' = q + f$
 - Every two byz-quorums intersect with $f+1$ acceptors
 - $f+1$ acceptors must contain a real acceptor.
- Simplify the problem: assume ballot-leaders cannot lie

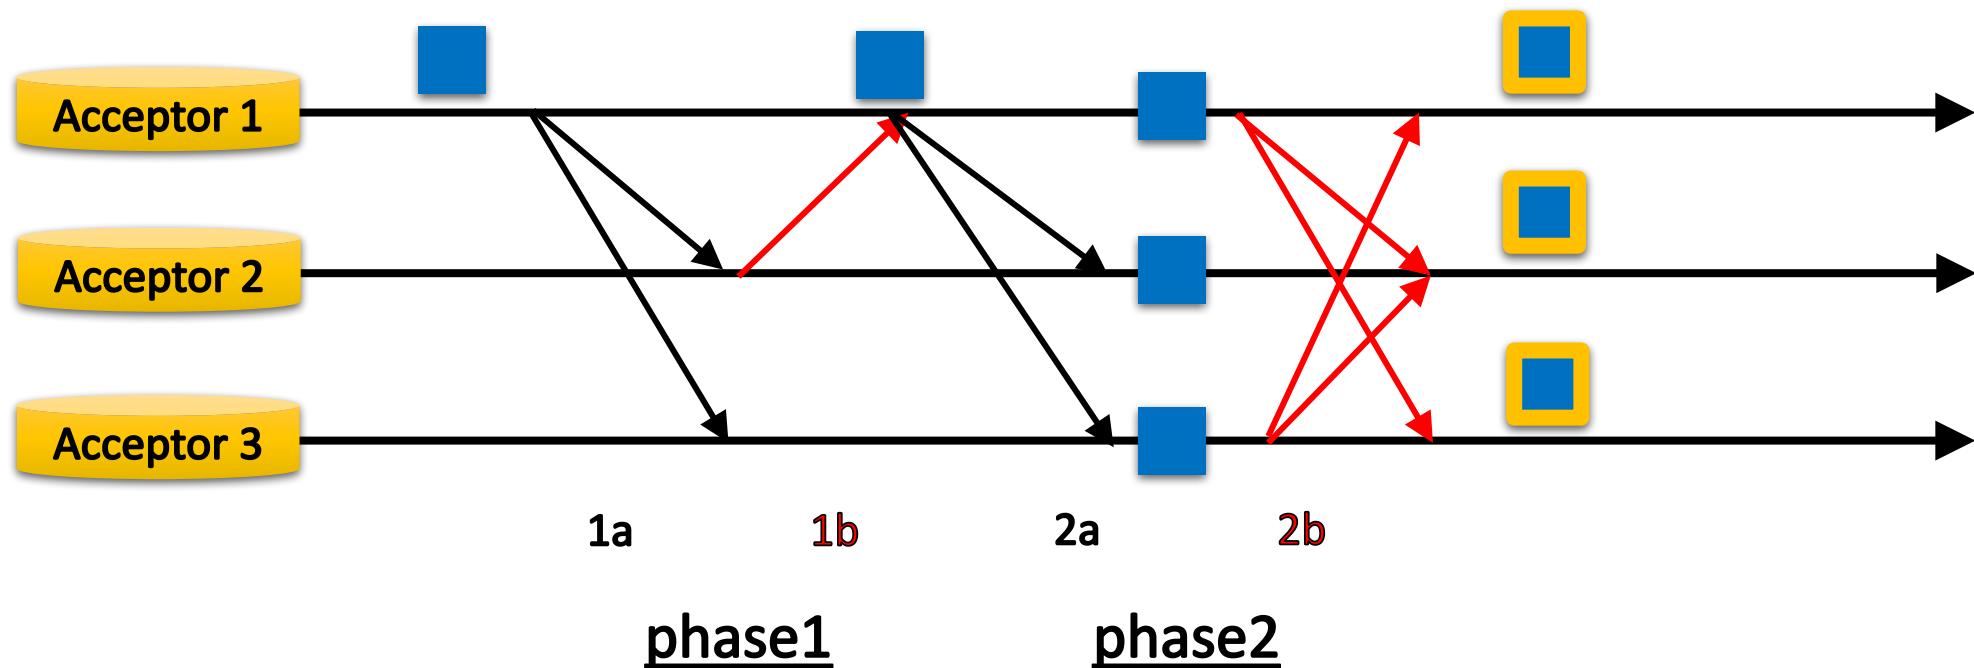


$N=3, q=2, \text{intersection}=1$

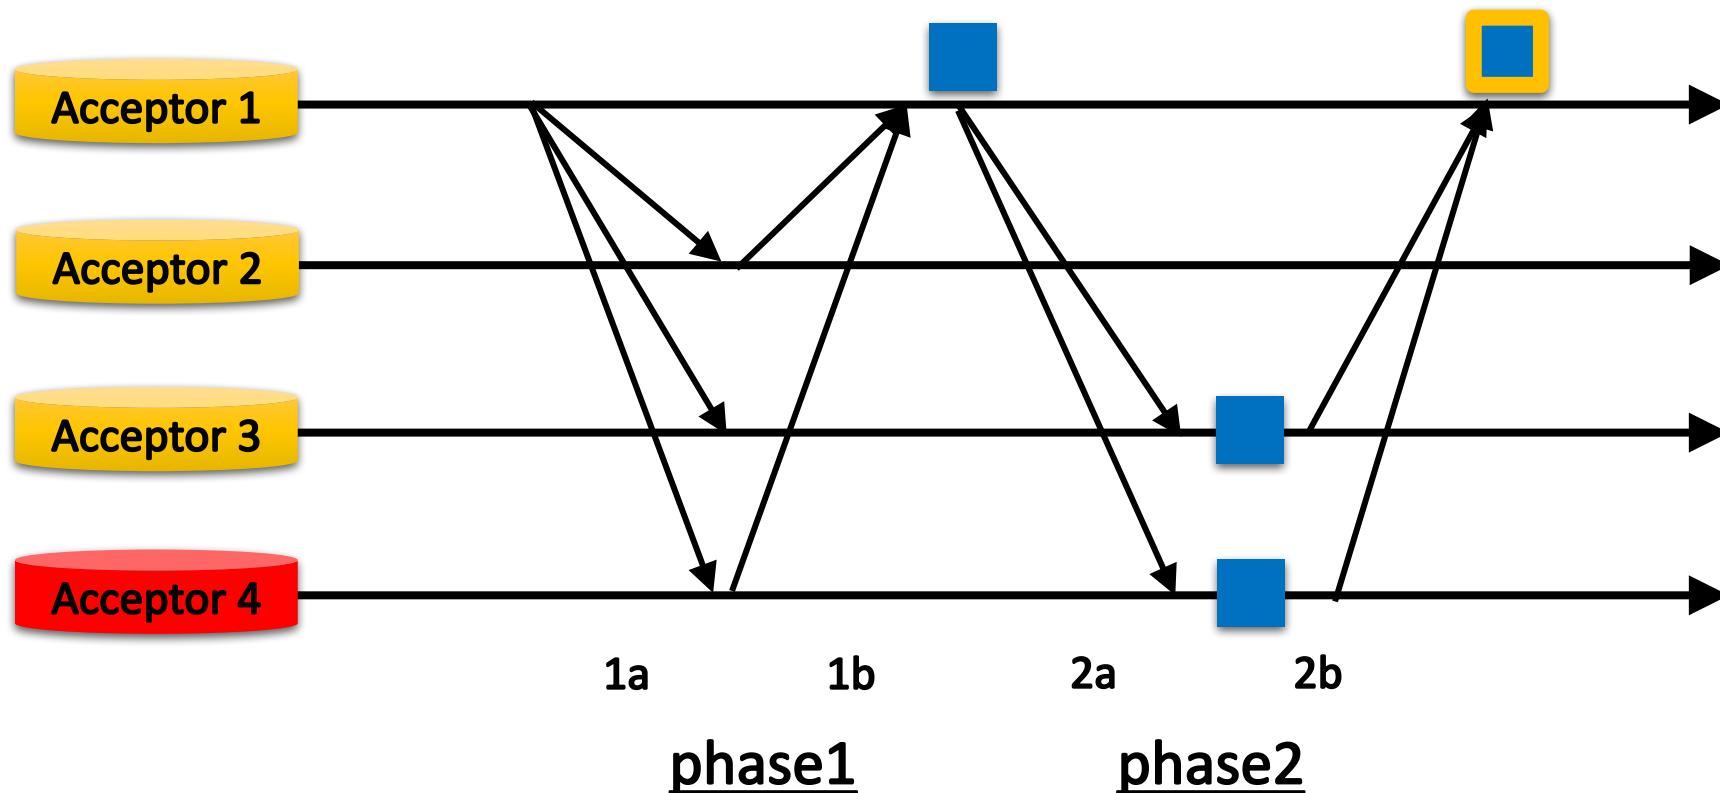
$f=1, N'=4, q'=3, \text{intersection}=2$

Liars in the system

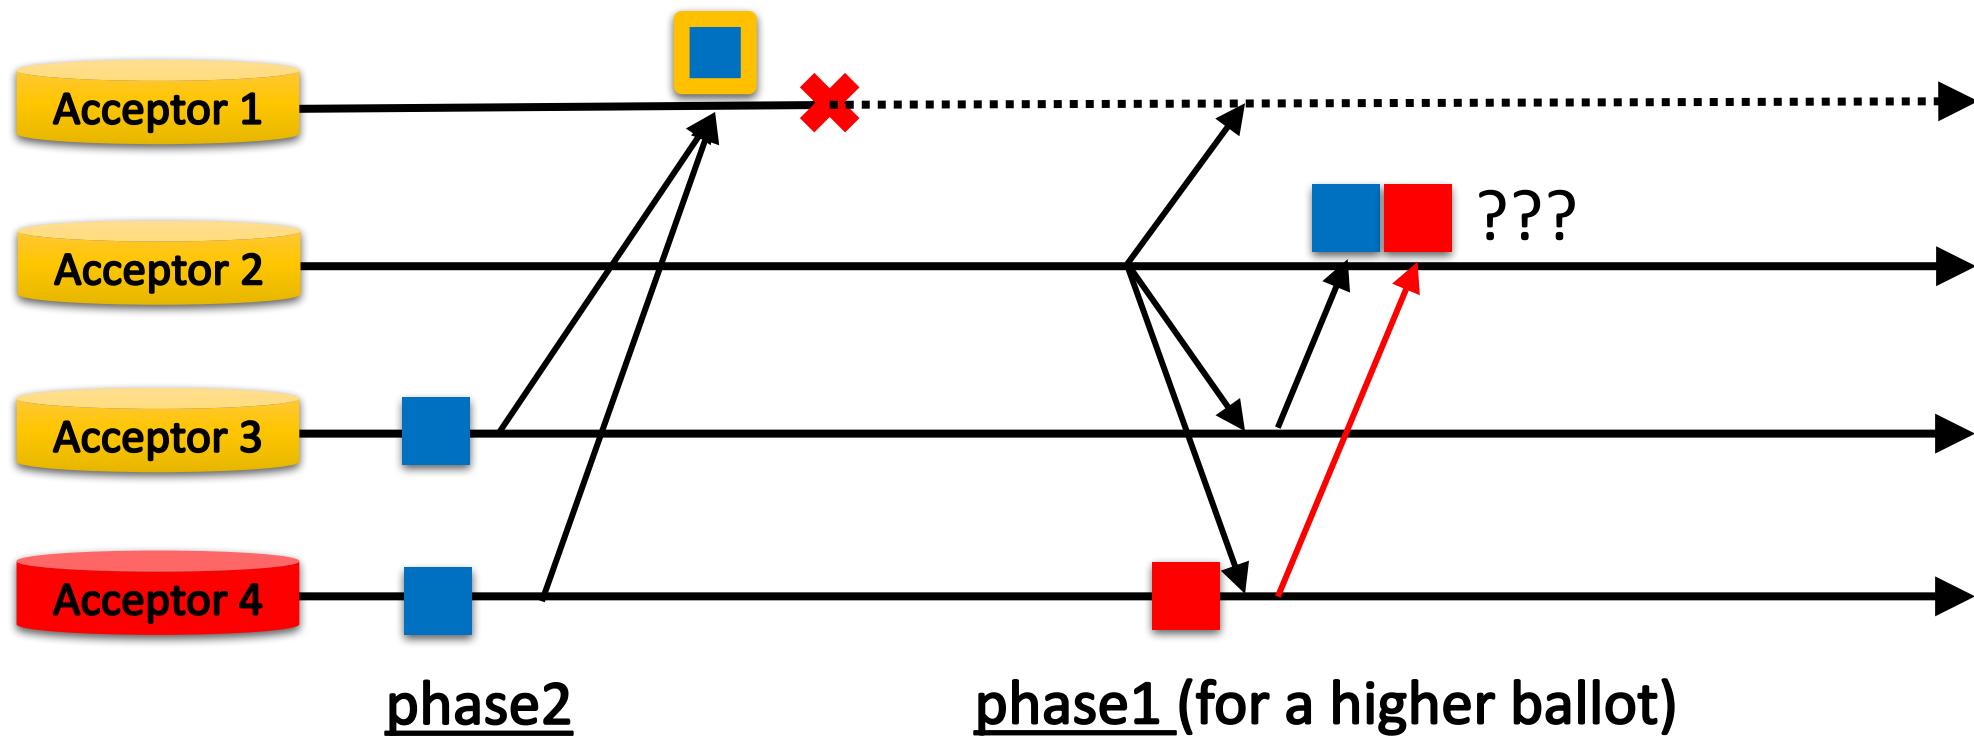
- Lies in Phase 1a, 2a
- Lies in Phase 1b, 2b



Modified Protocols with byz-quorum

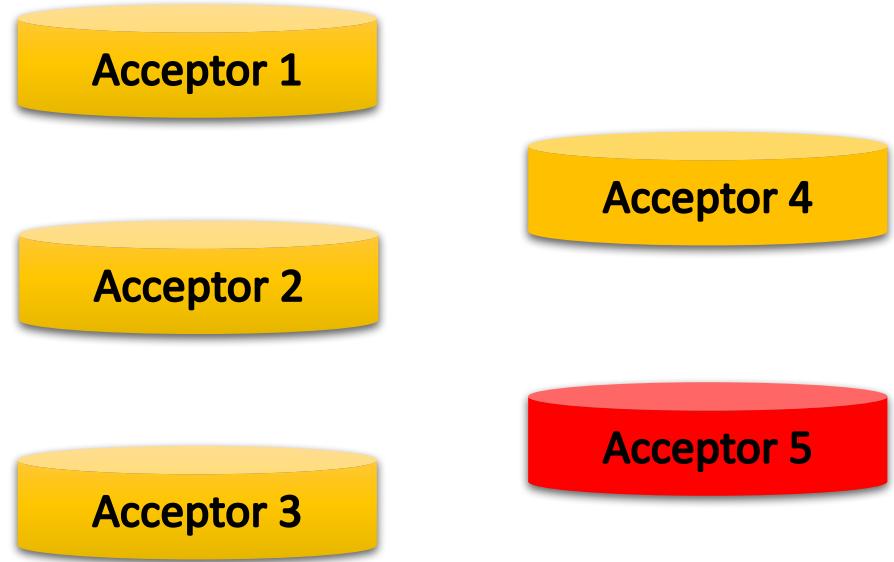


Acceptor can lie in phase 1b



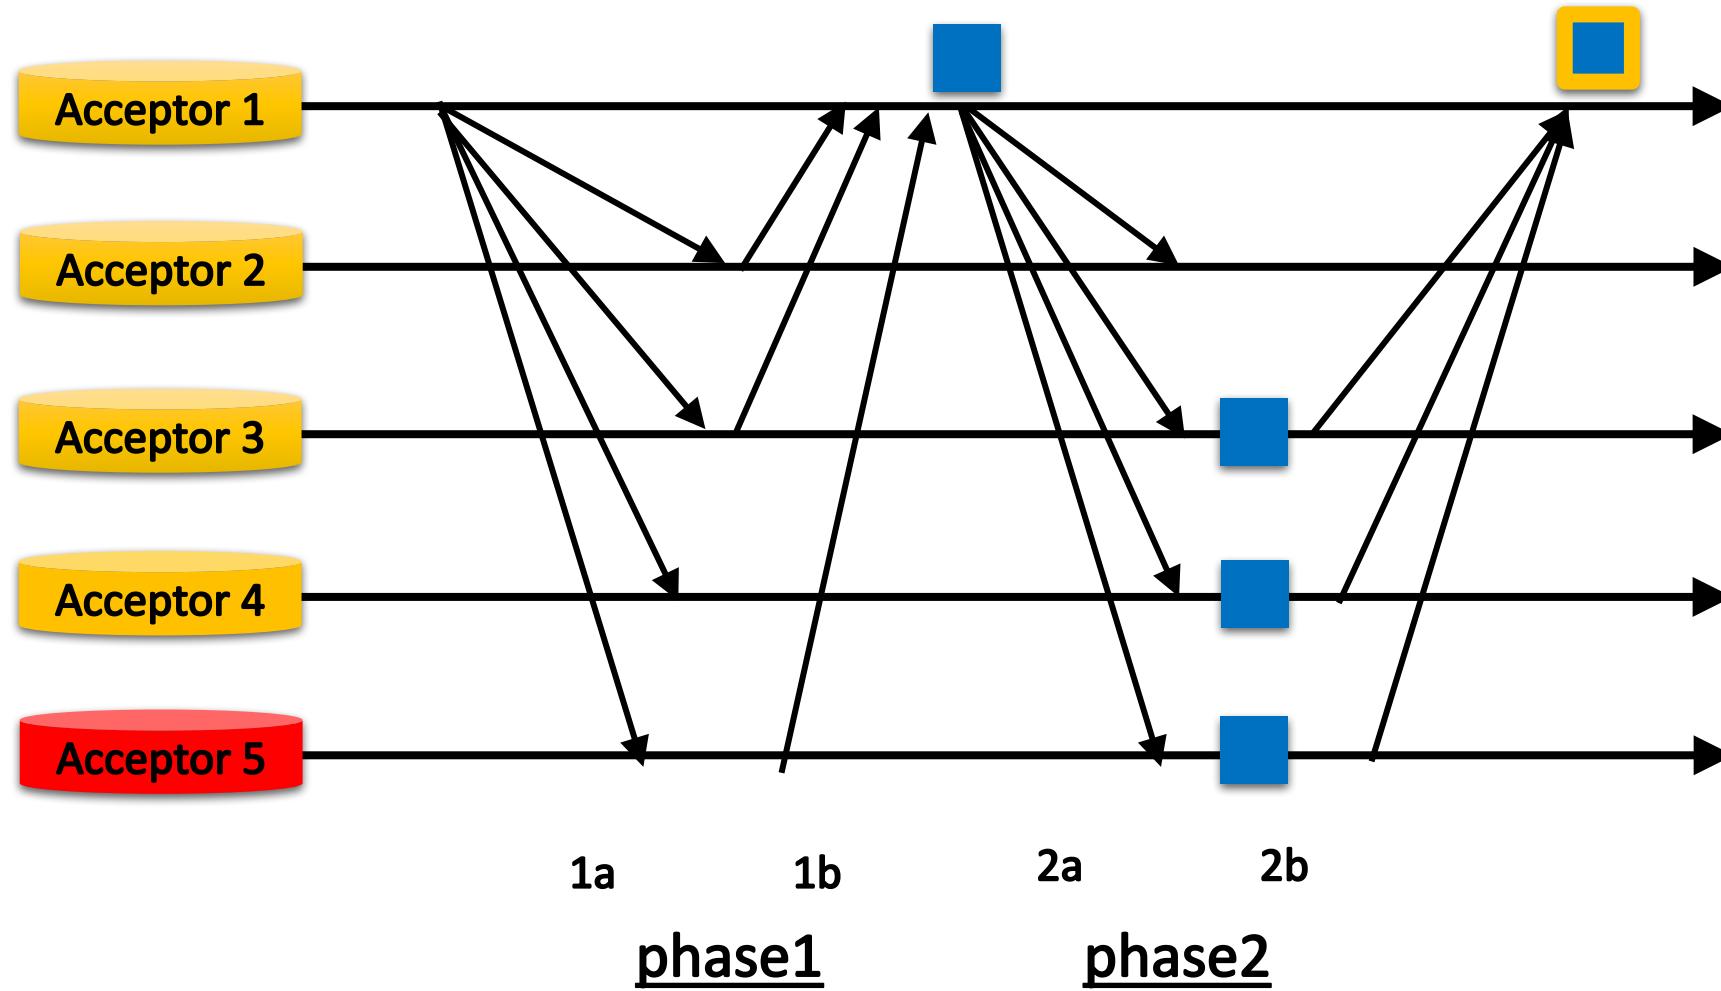
How to outvote liar

- A possible solution: increase the quorum size
 - Original quorum size $q = 2f+1$,
 - Original quorum Intersection = $f+1$
 - Original $N = 3f+1$
 - Byz-quorum $q' = 3f+1$
 - Byz-quorum intersection = $2f+1$
 - Byz acceptors $N' = 4f+1$

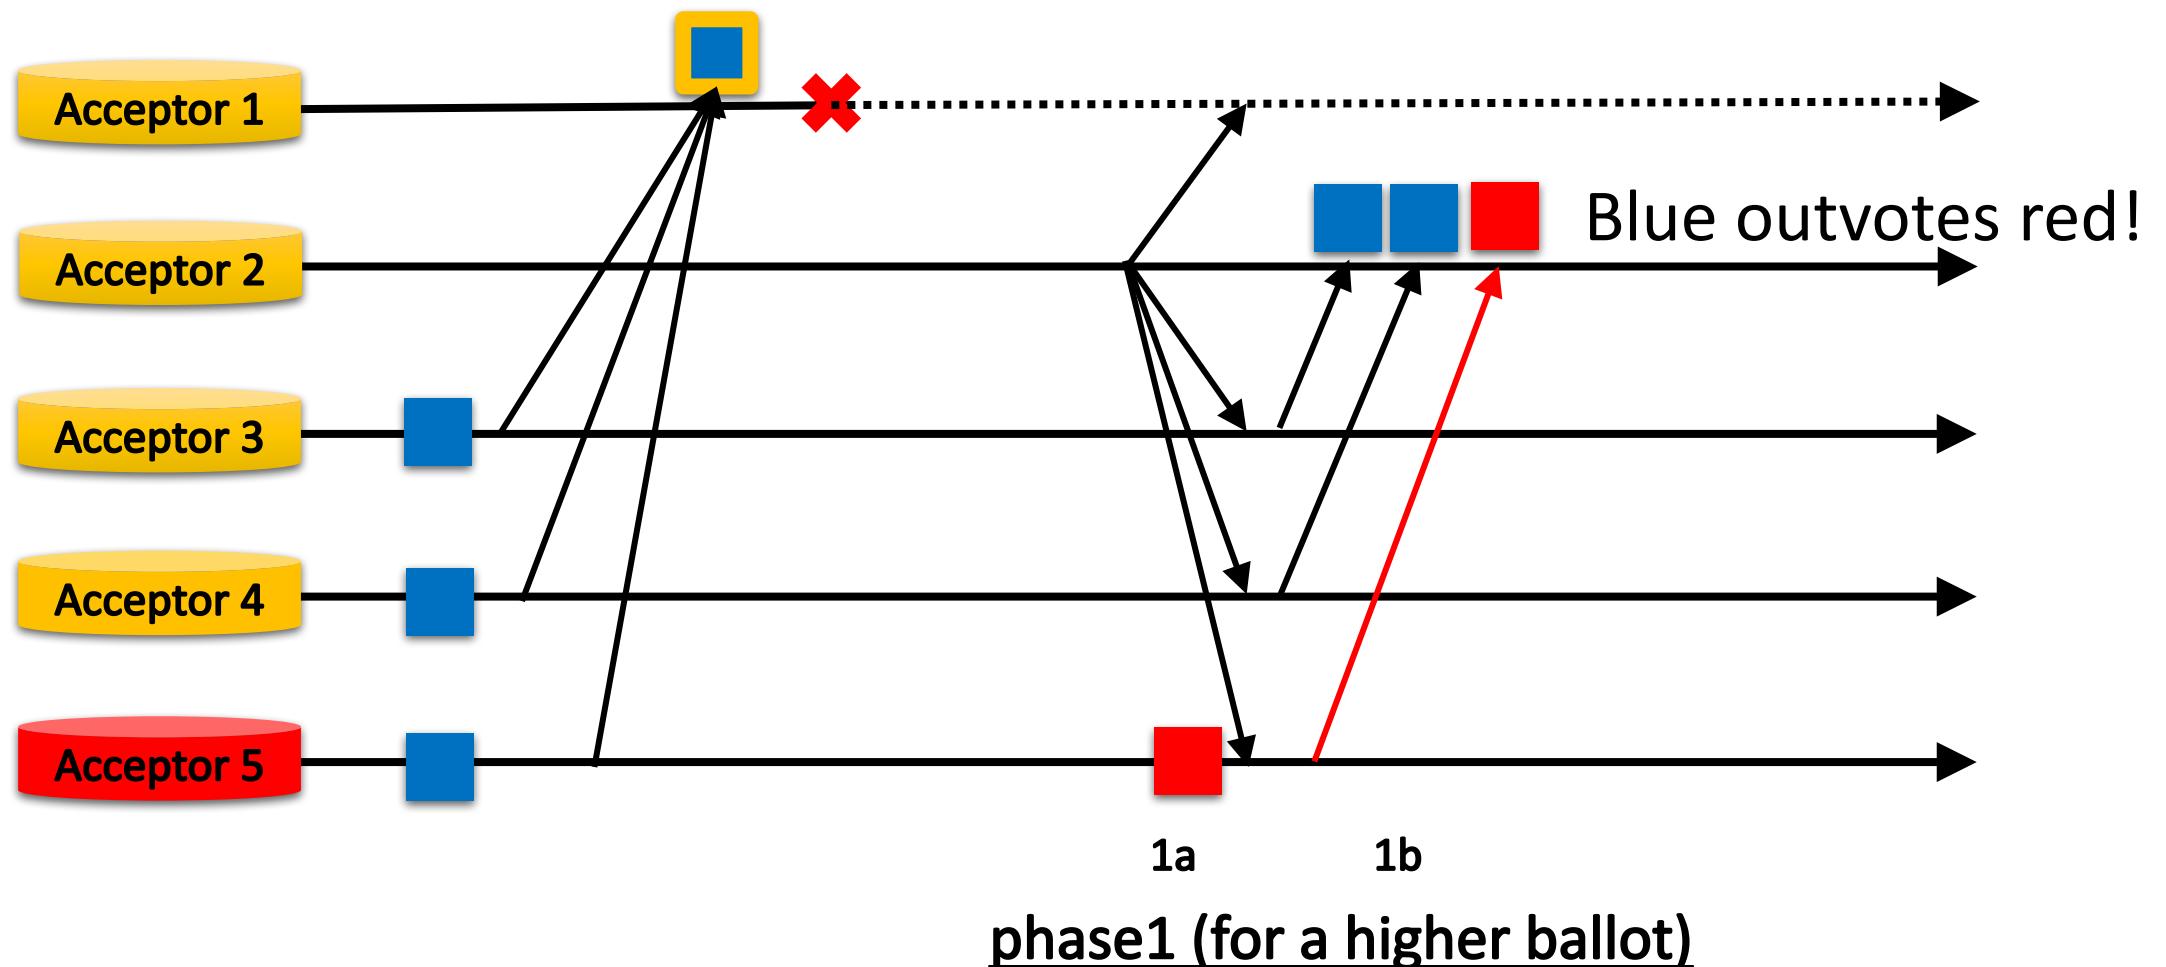


$N=4, q=3, \text{intersection}=2$
 $f=1, N'=5, q'=4, \text{intersection}=3$

Modified Protocols with byz-quorum



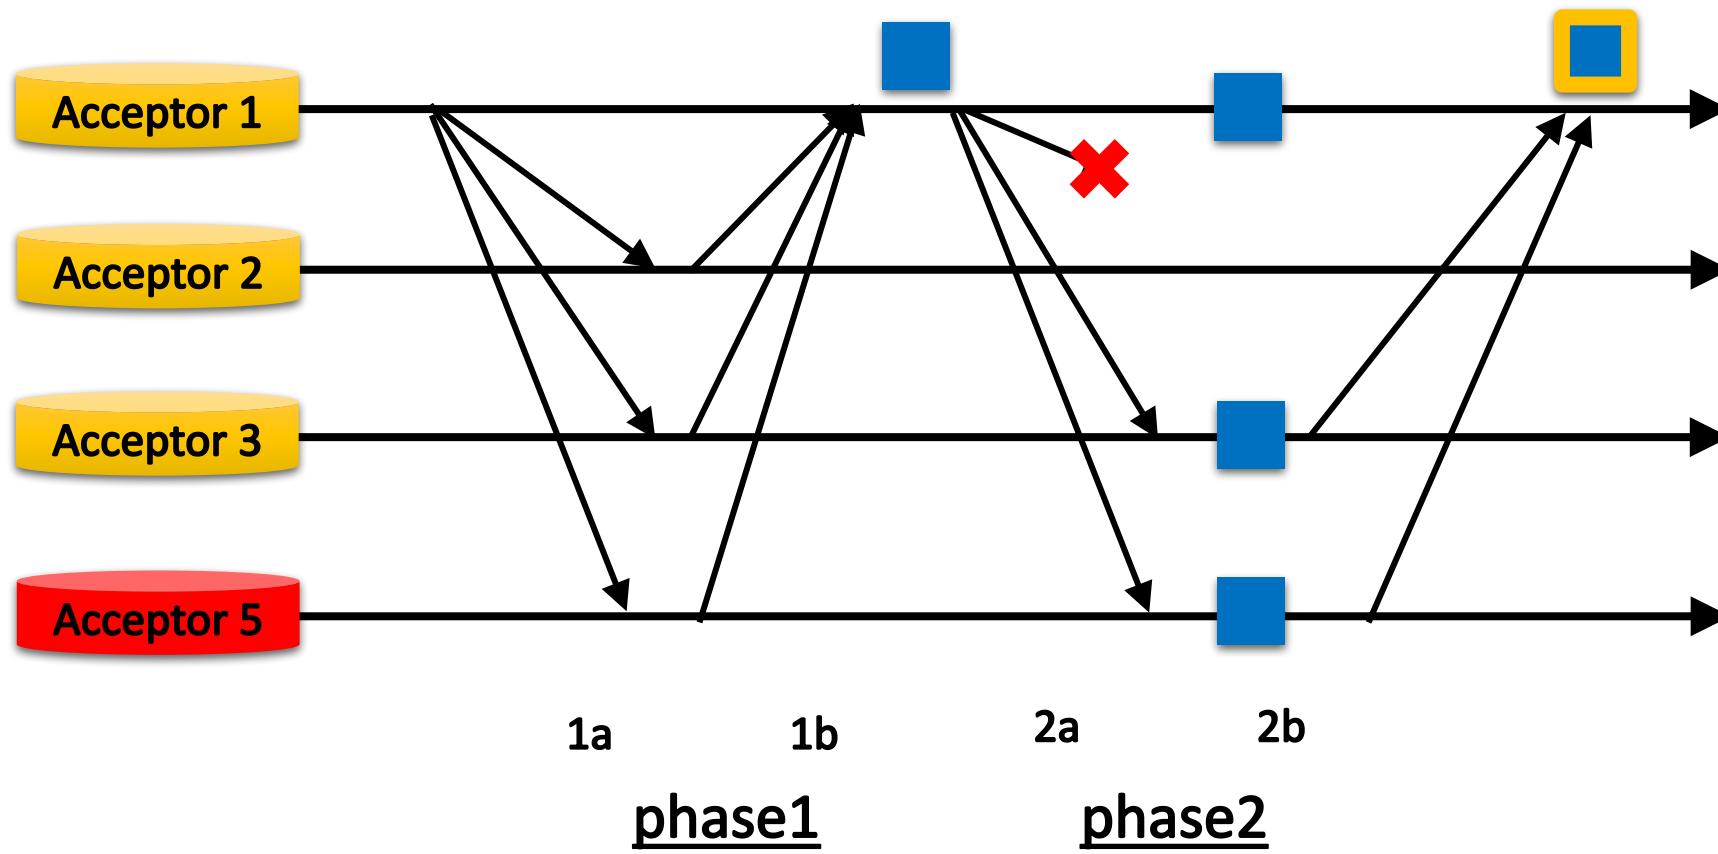
Example



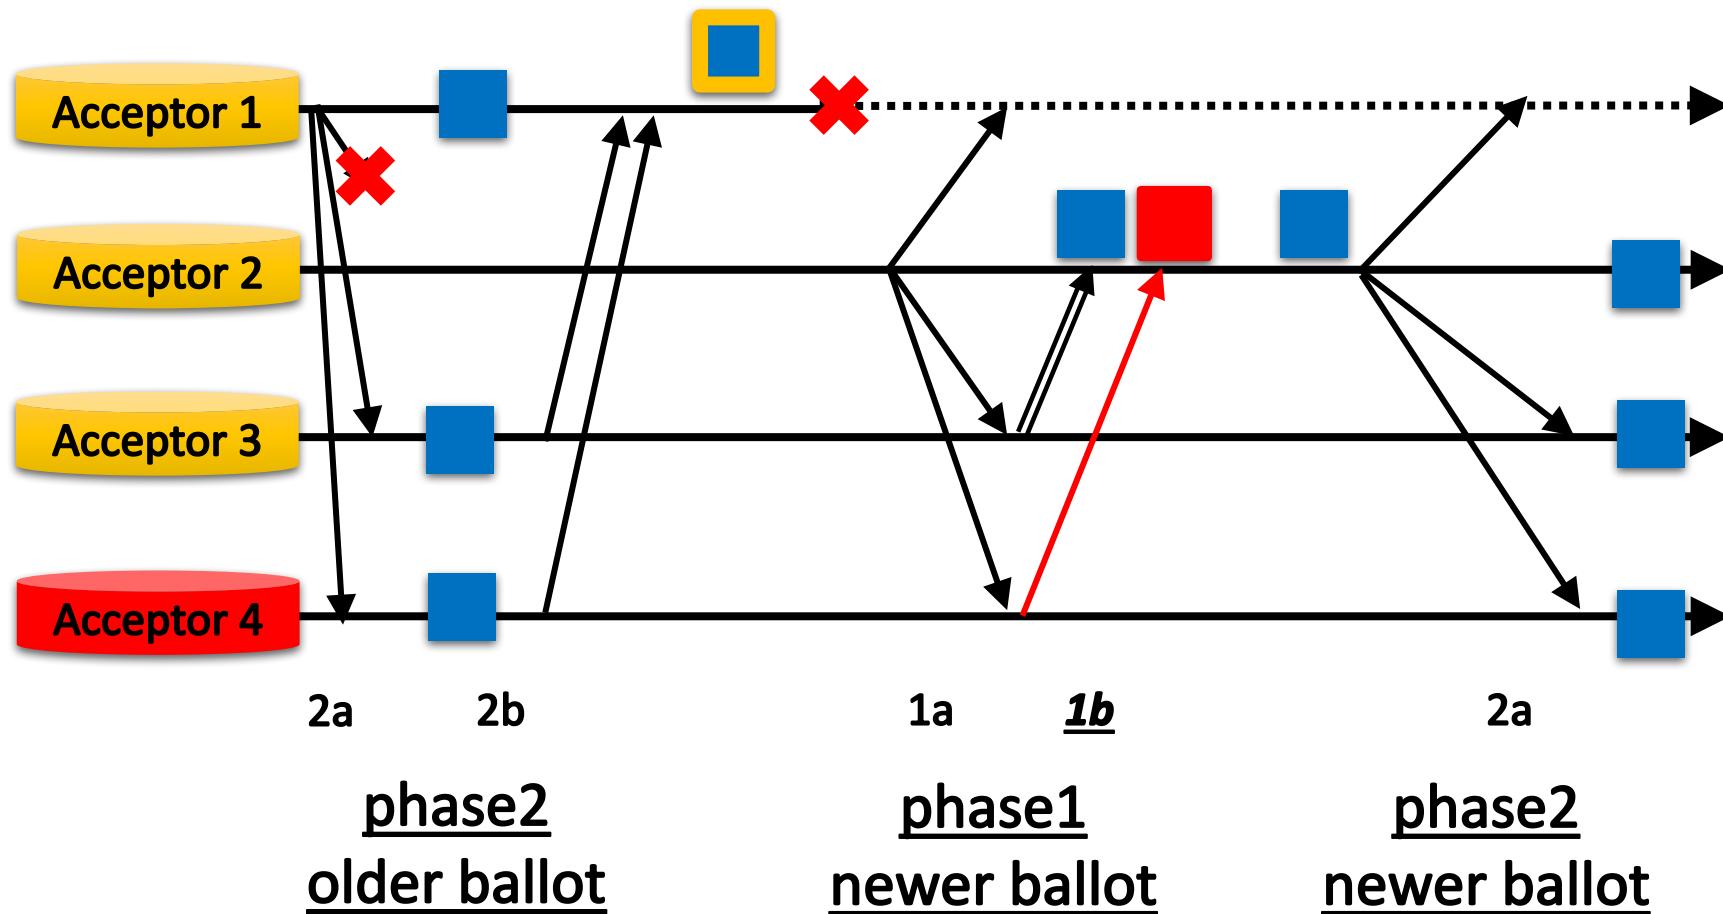
Problem now: too many acceptors

- $N' = N + f = 4f + 1$
- Can we shrink the size of acceptors?
 - From $4f+1$ to $3f+1$?
- Alternative solution: disable the lies in phase 1b
 - Require the value in 1b message be attached with a “proof”
 - Assert that this value was actually a safe value, not a fake value
 - The proof is implied by phase 2a message (assumes no lies of 2a)

Example



Example (cont.)



Question

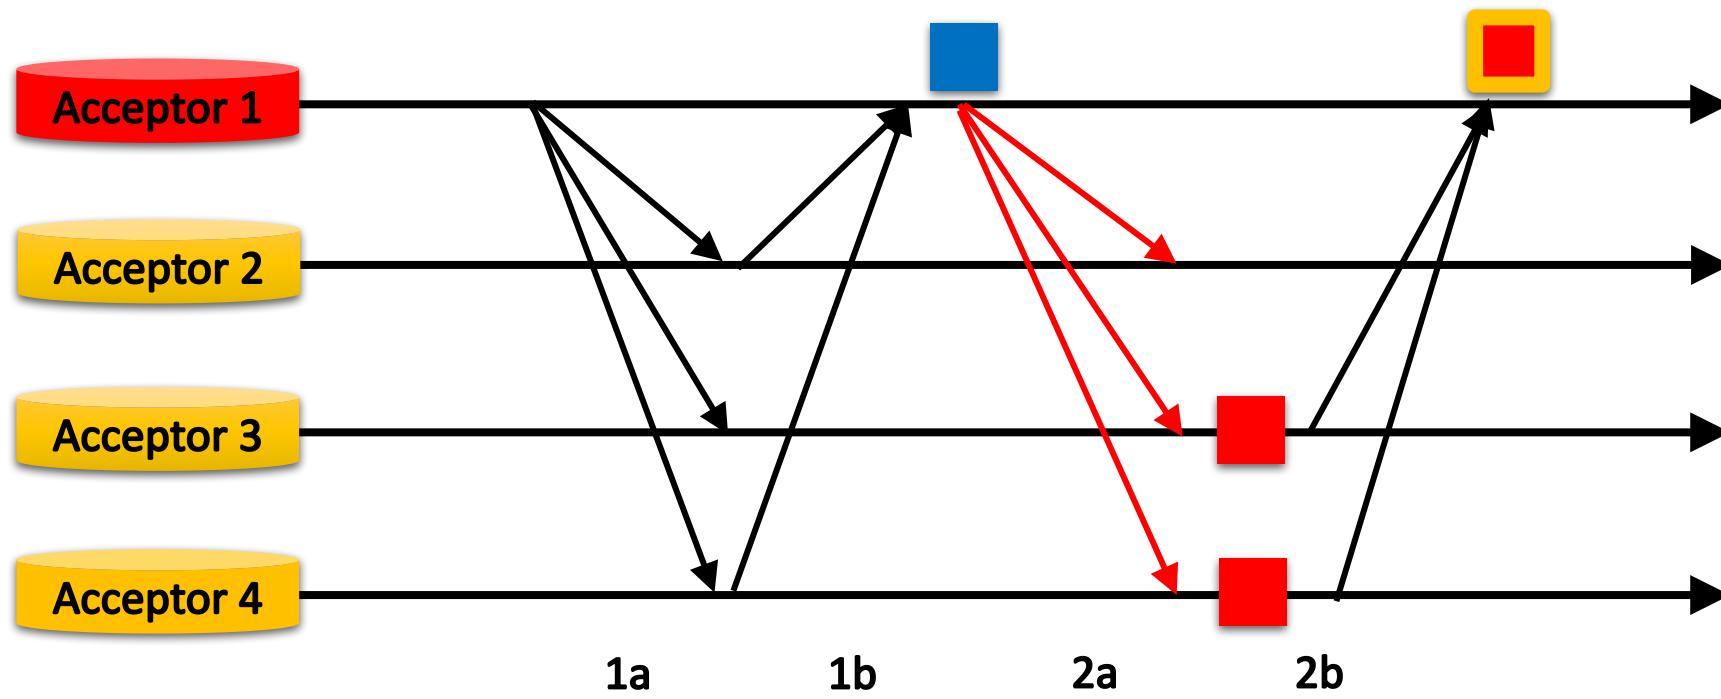
- In phase 1b, can acceptor still lie without being caught?
- If so, does it hurt?
- What about lies in phase 2b?

Structure

- Background
- Paxos Review
- ByzPaxos (Non-leader lies)
- **ByzPaxos (Leader lies)**
 - Break down phase 2a
 - Cooperative phase 2a
 - Multi-instance

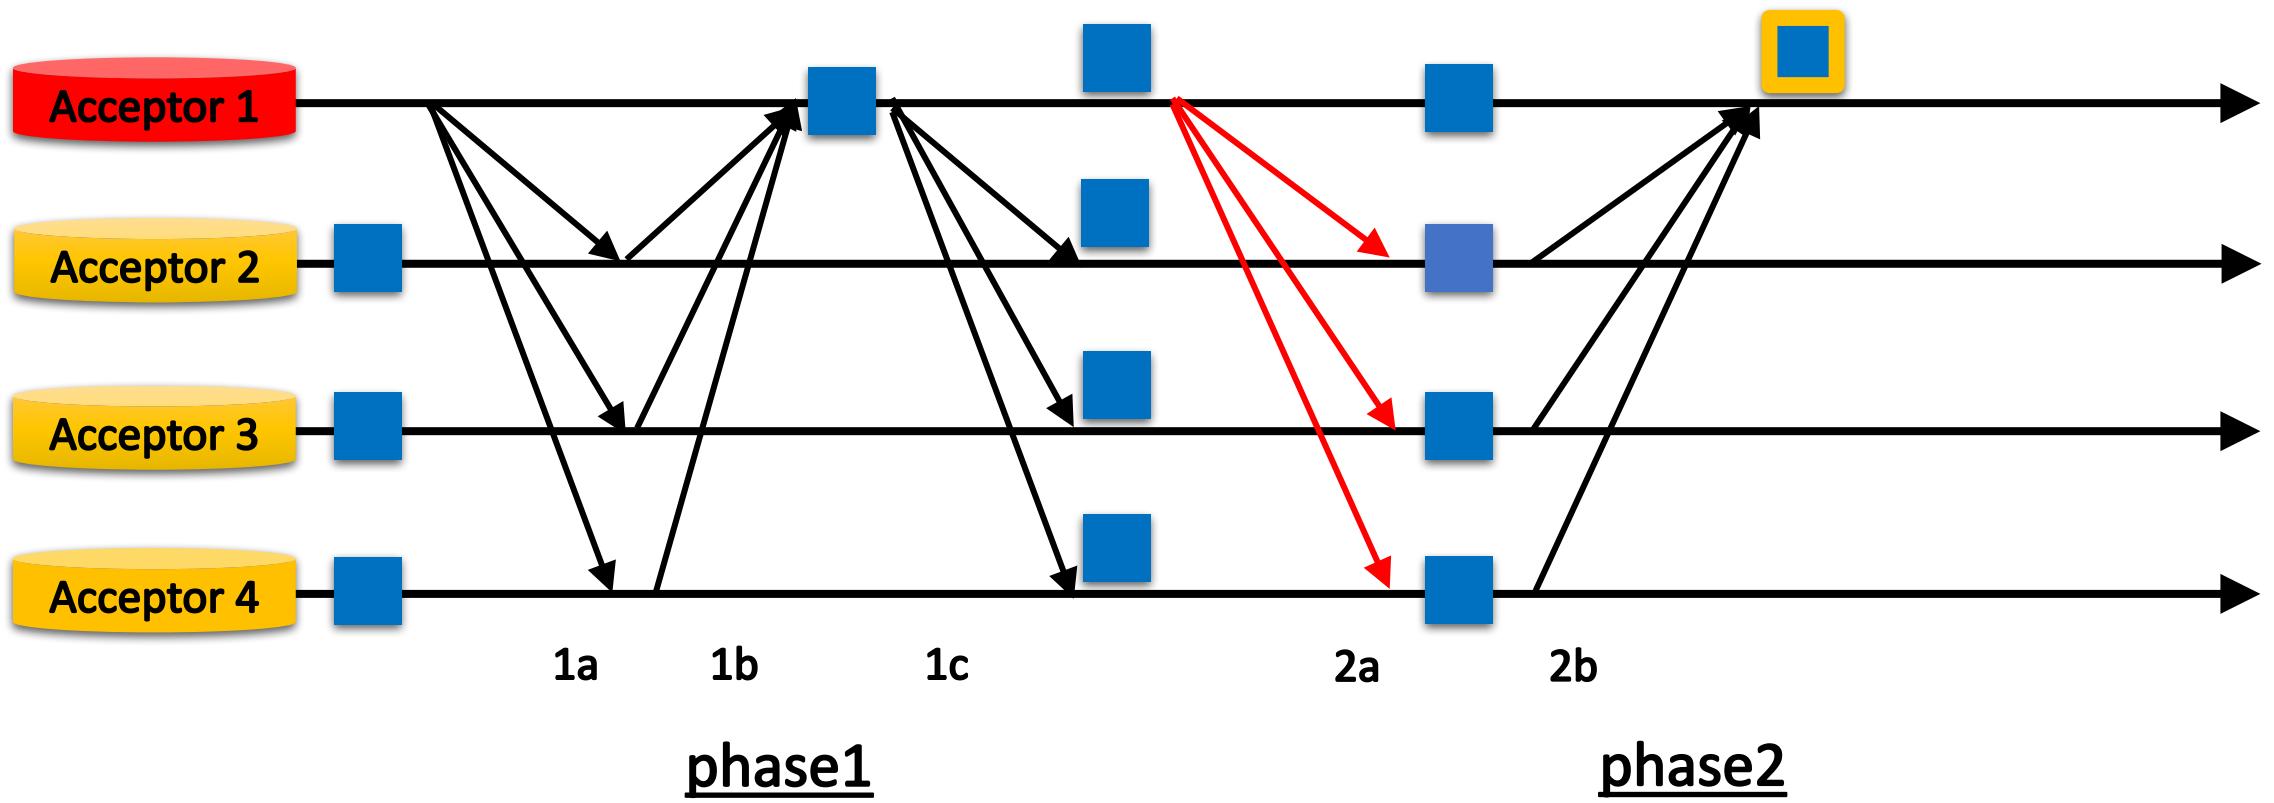
A fake acceptor can lie as leader for a ballot

- Cannot lie in phase 1a
- Lie in phase 2a matters!
 - Lying leader can feed any value
 - Acceptor does not know how to validate if it is “legal” (safe)

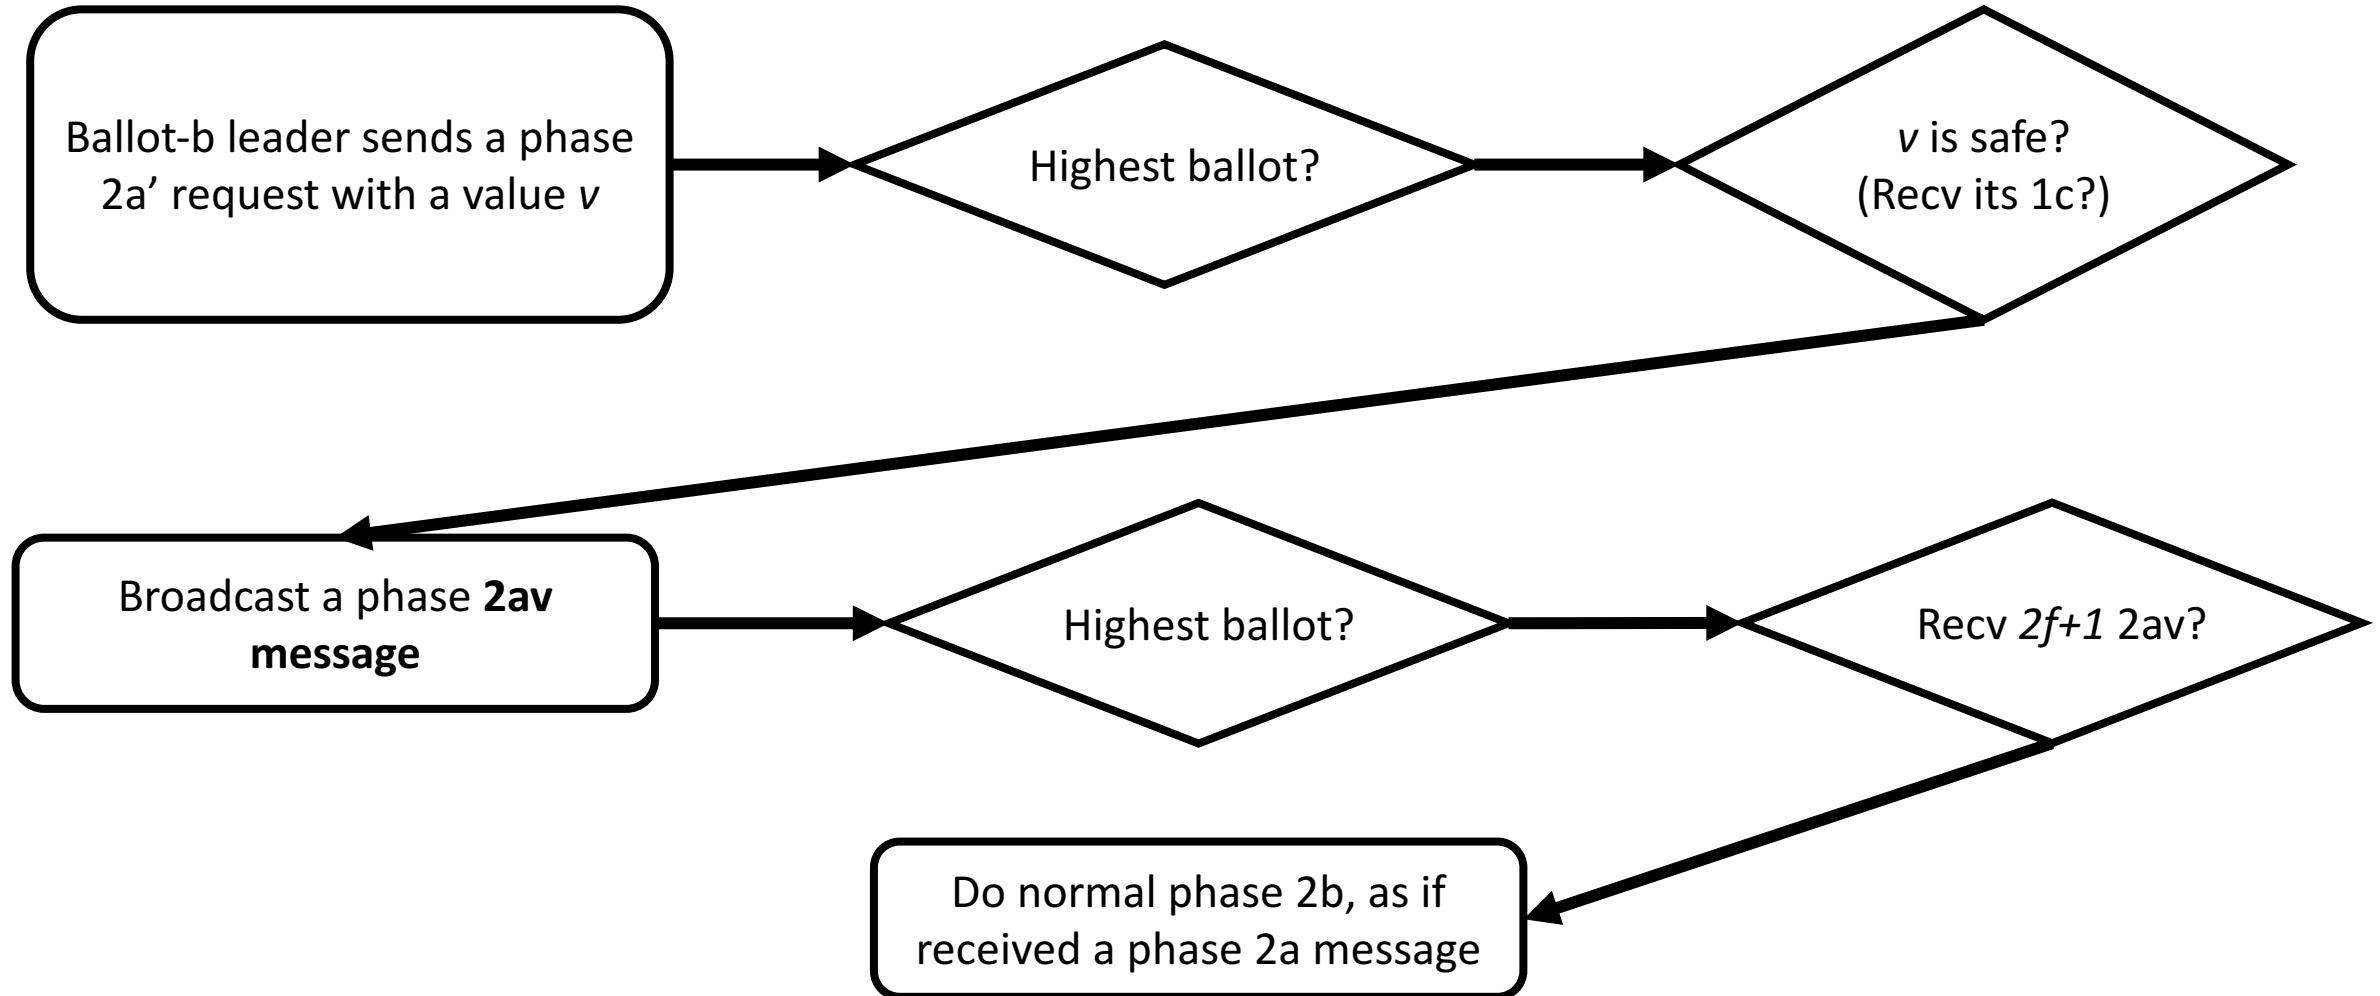


Extract phase 1c from phase 2a

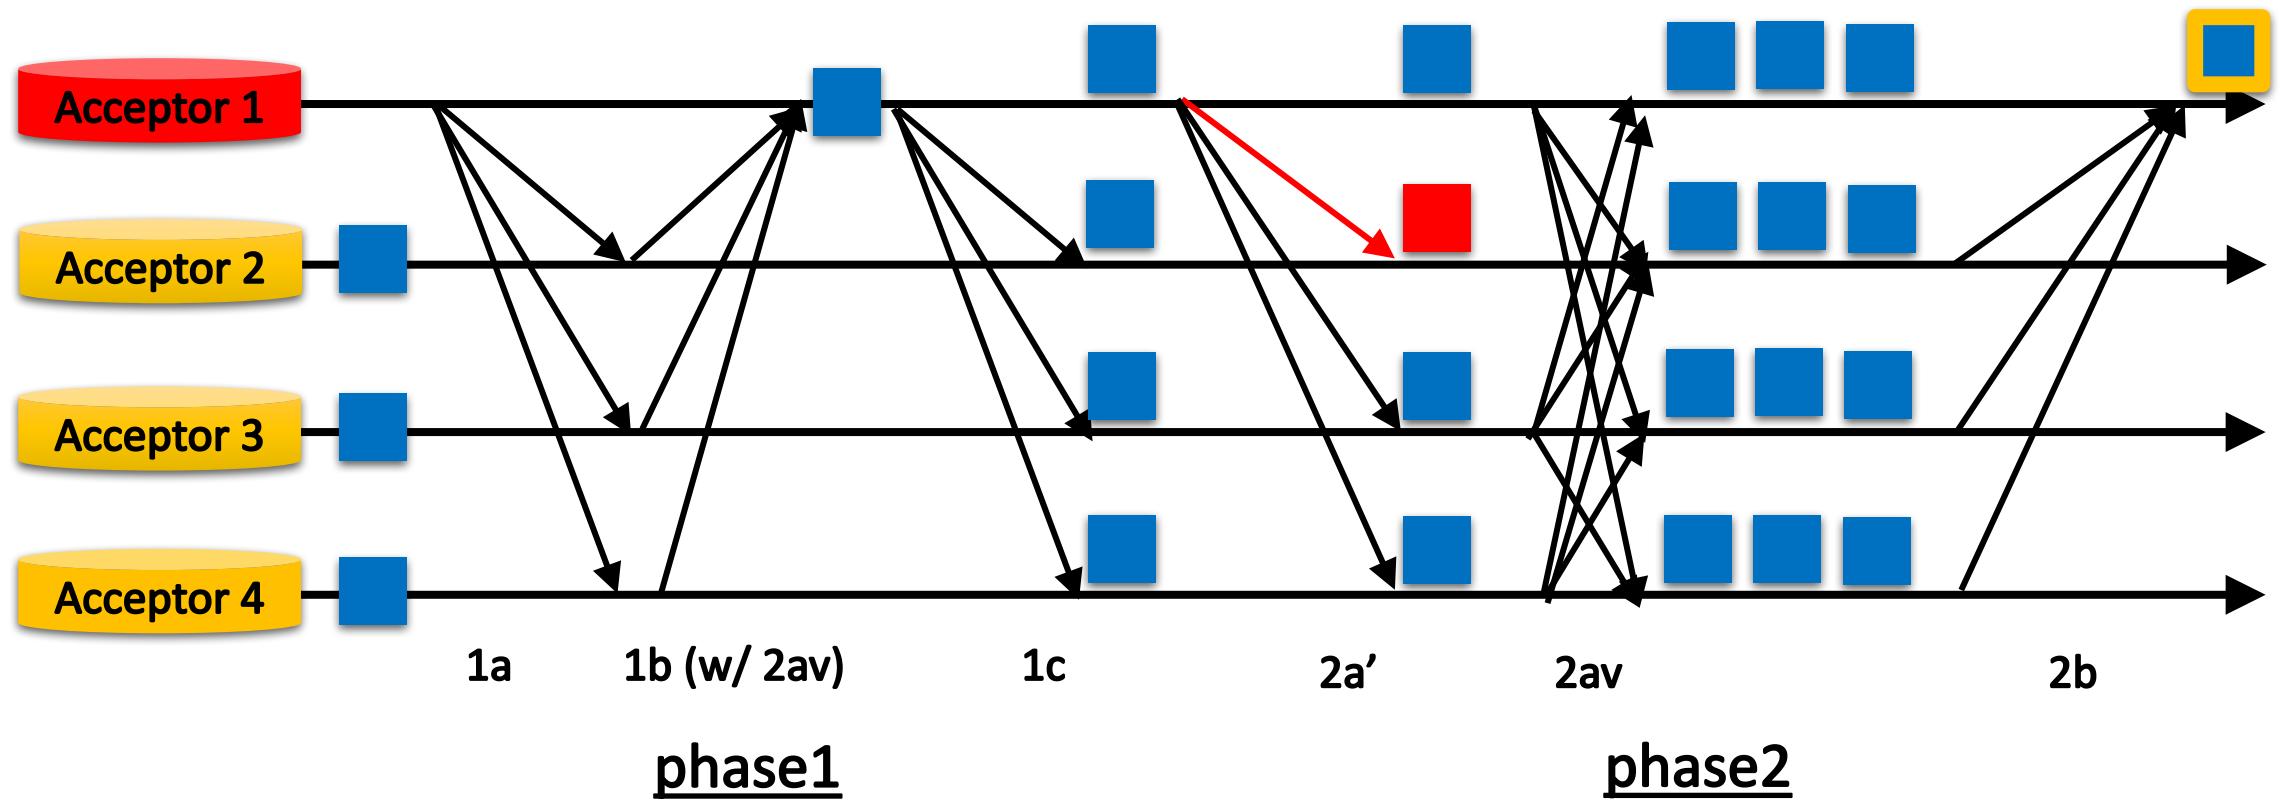
- Give acceptors more information to know what is safe
- Phase 1c message for ballot b asserts some (all) values are safe



Dealing with lies w/ cooperative phase 2a

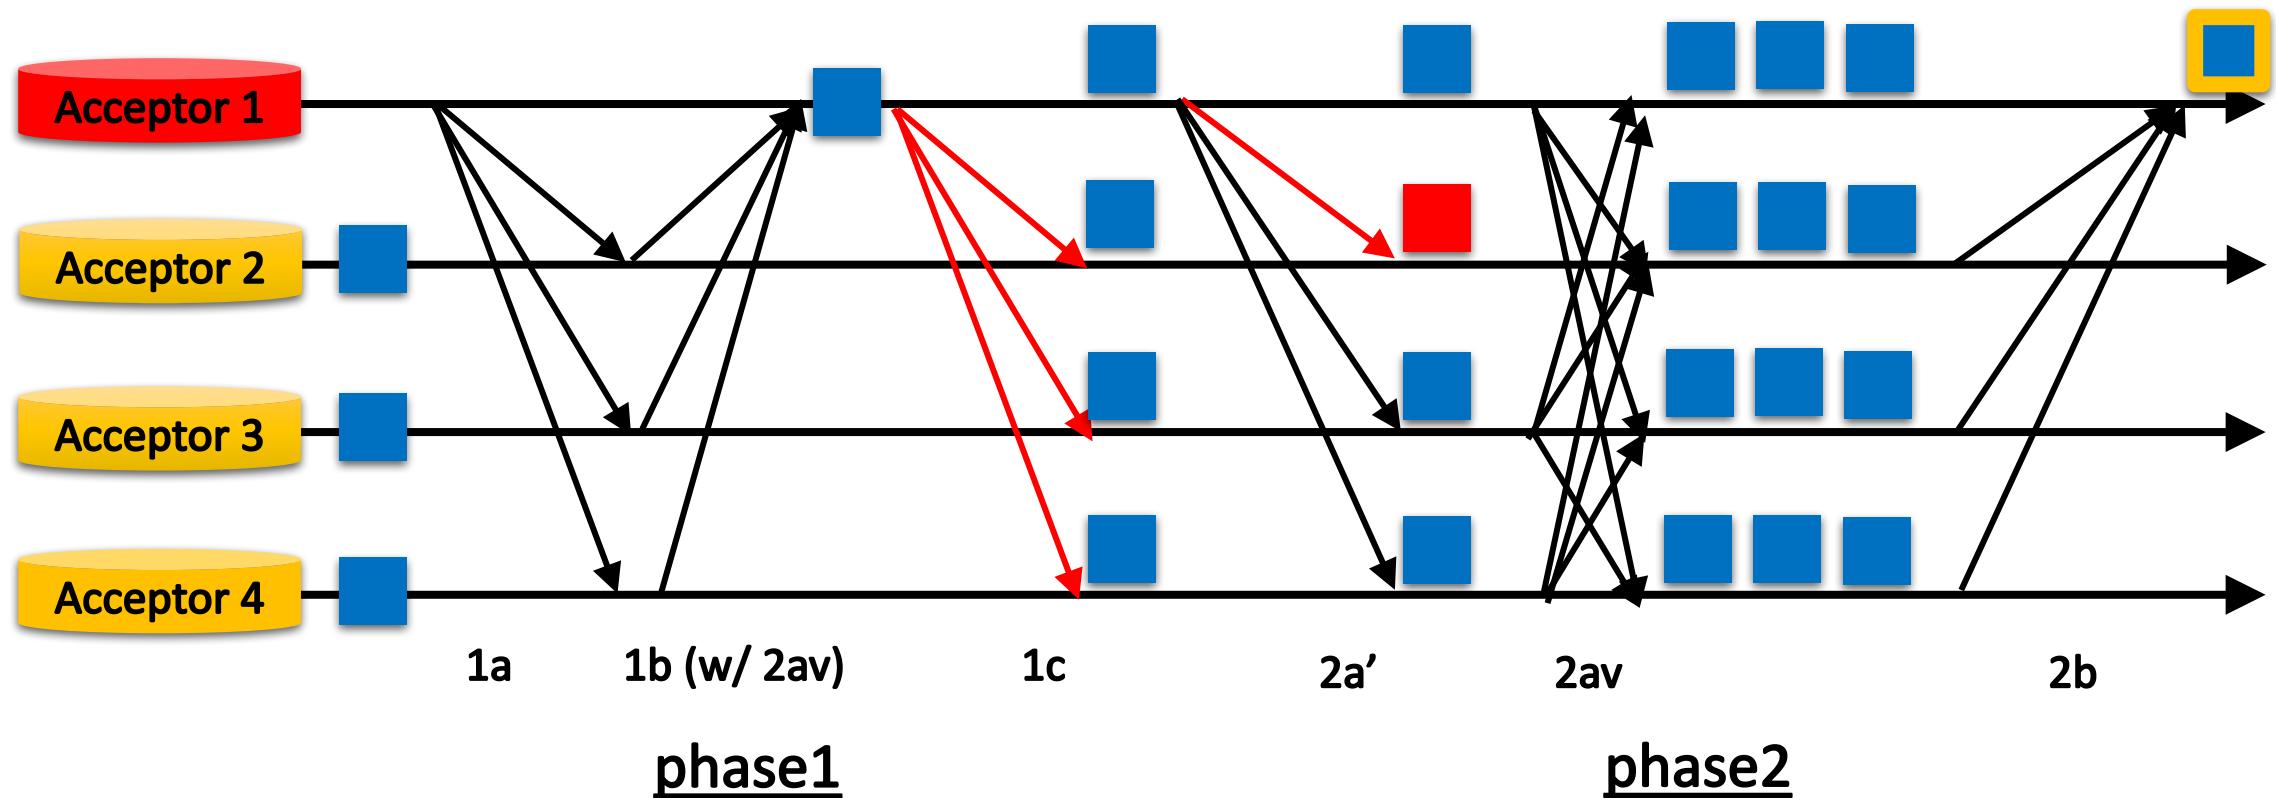


Example



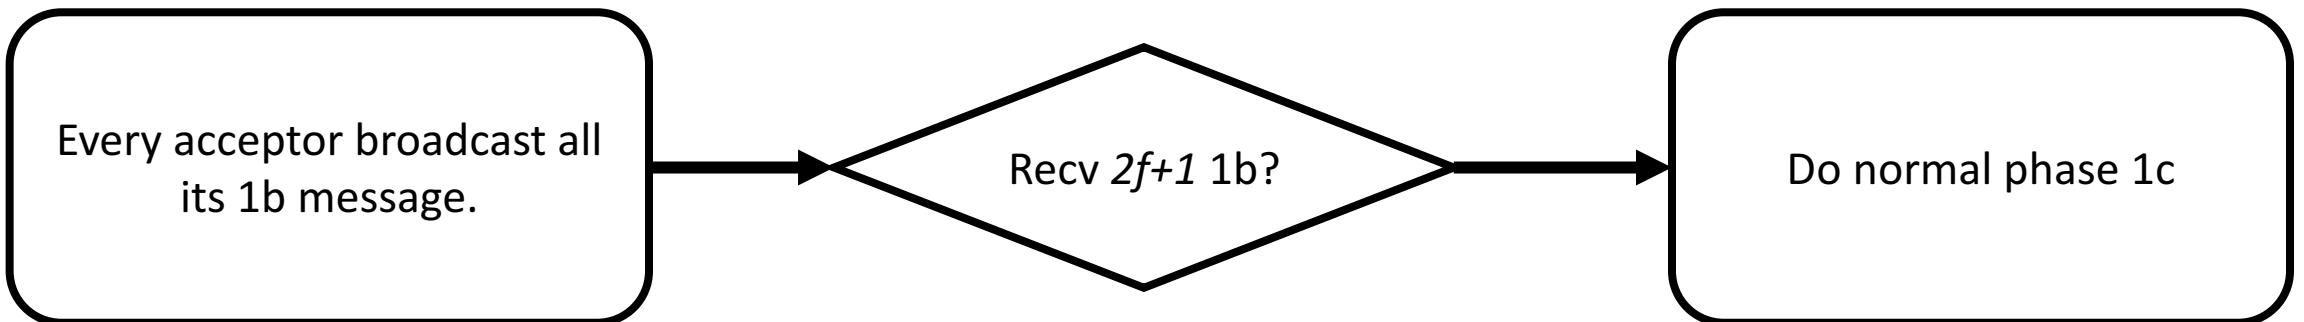
What about lies in phase 1c?

- Can we validate phase 1c message?

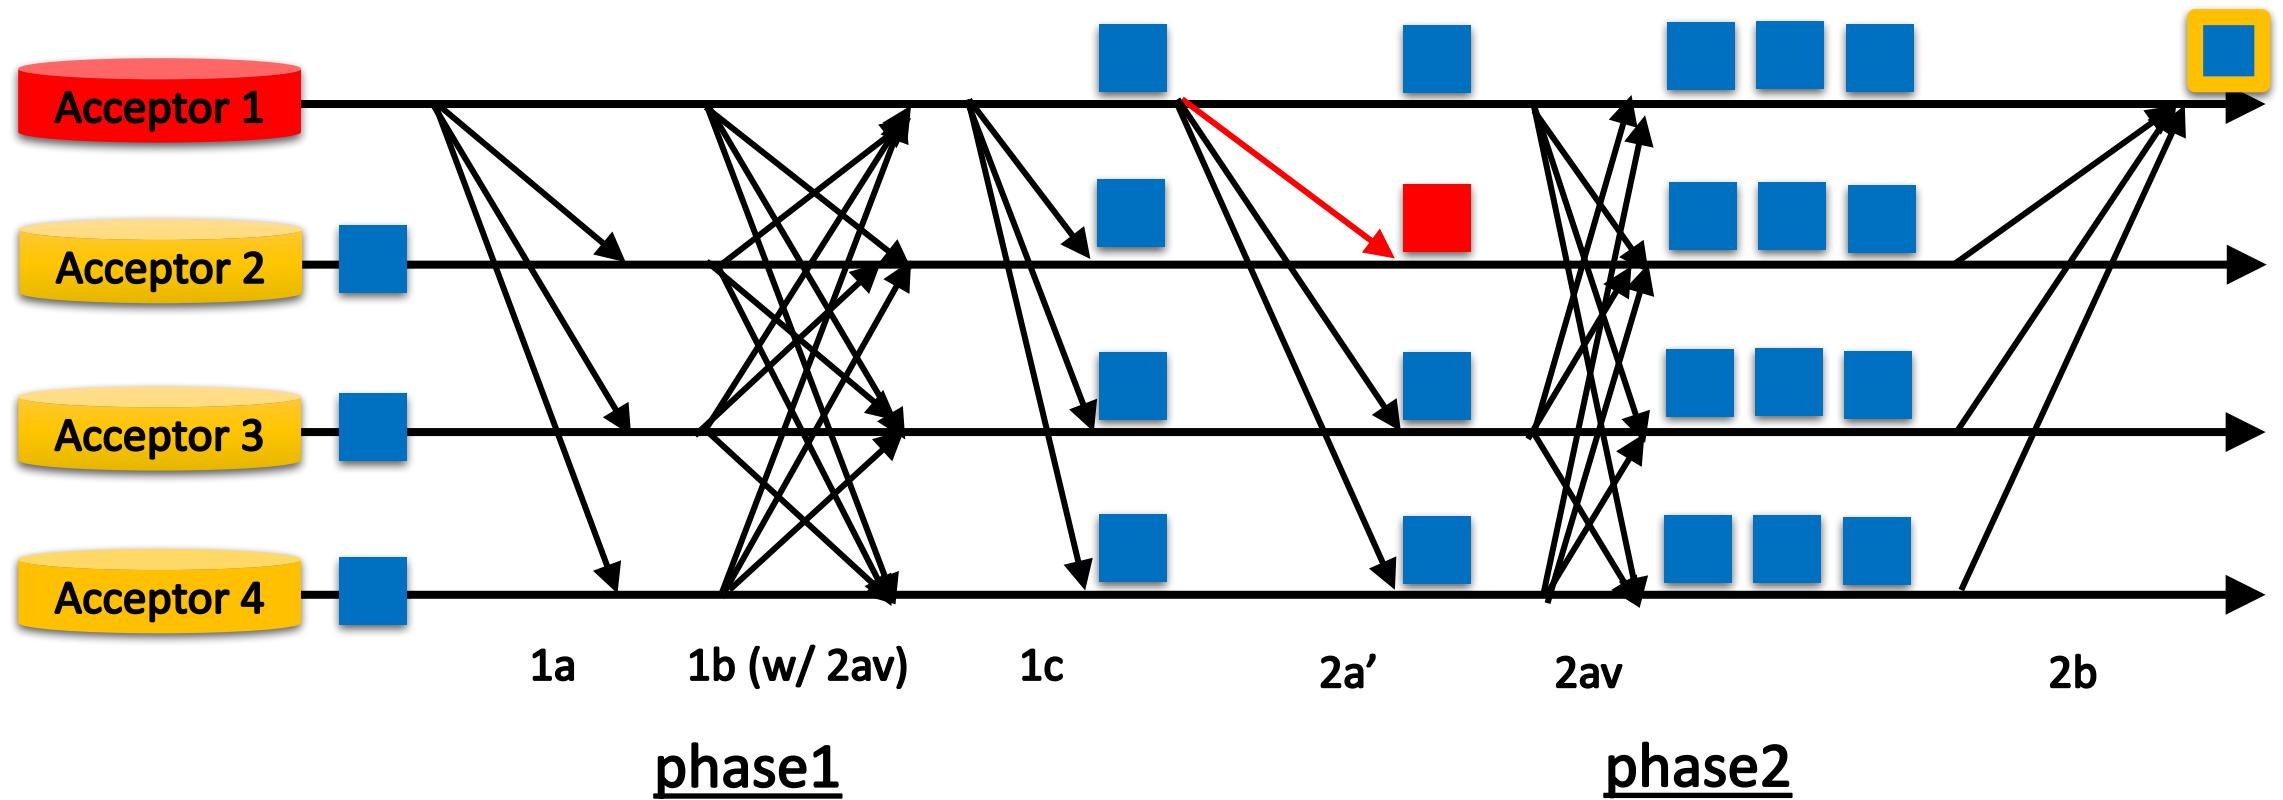


Deduct phase 1c from phase 1b

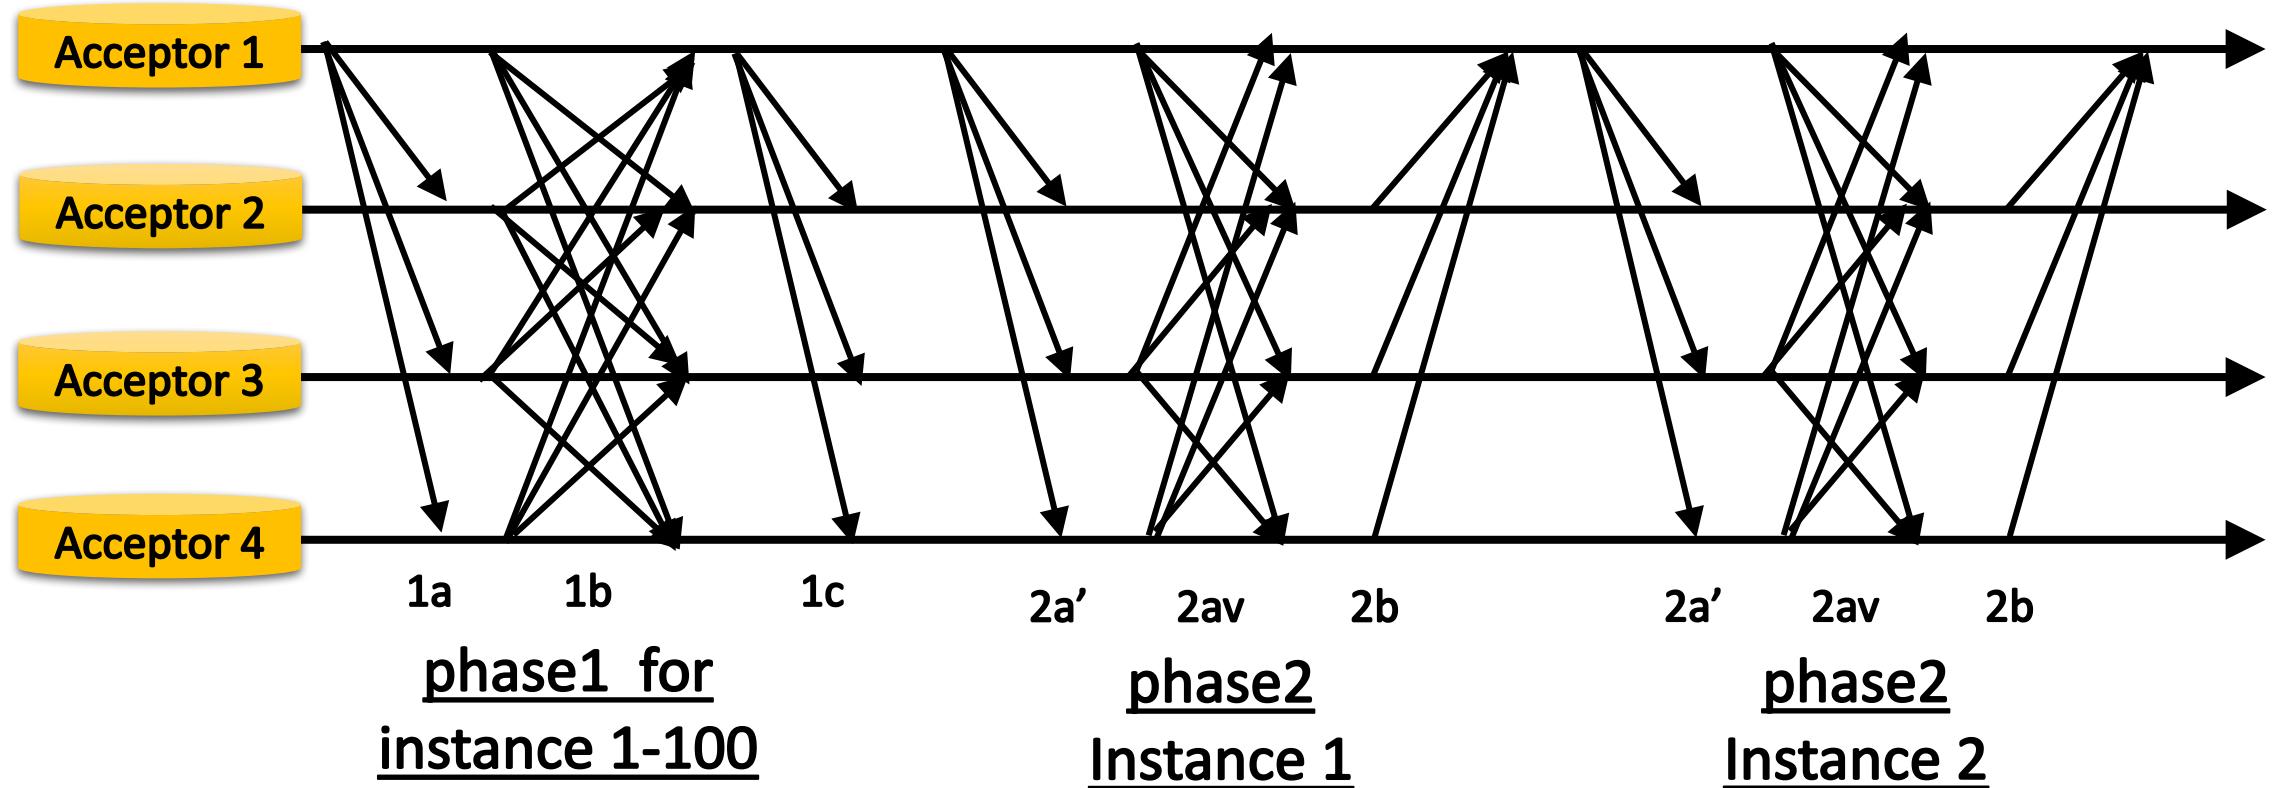
- The ballot- b leader sends 1c on receiving a byz-quorum of 1b
- Prevent the leader lying by making every acceptor repeat the process



Example

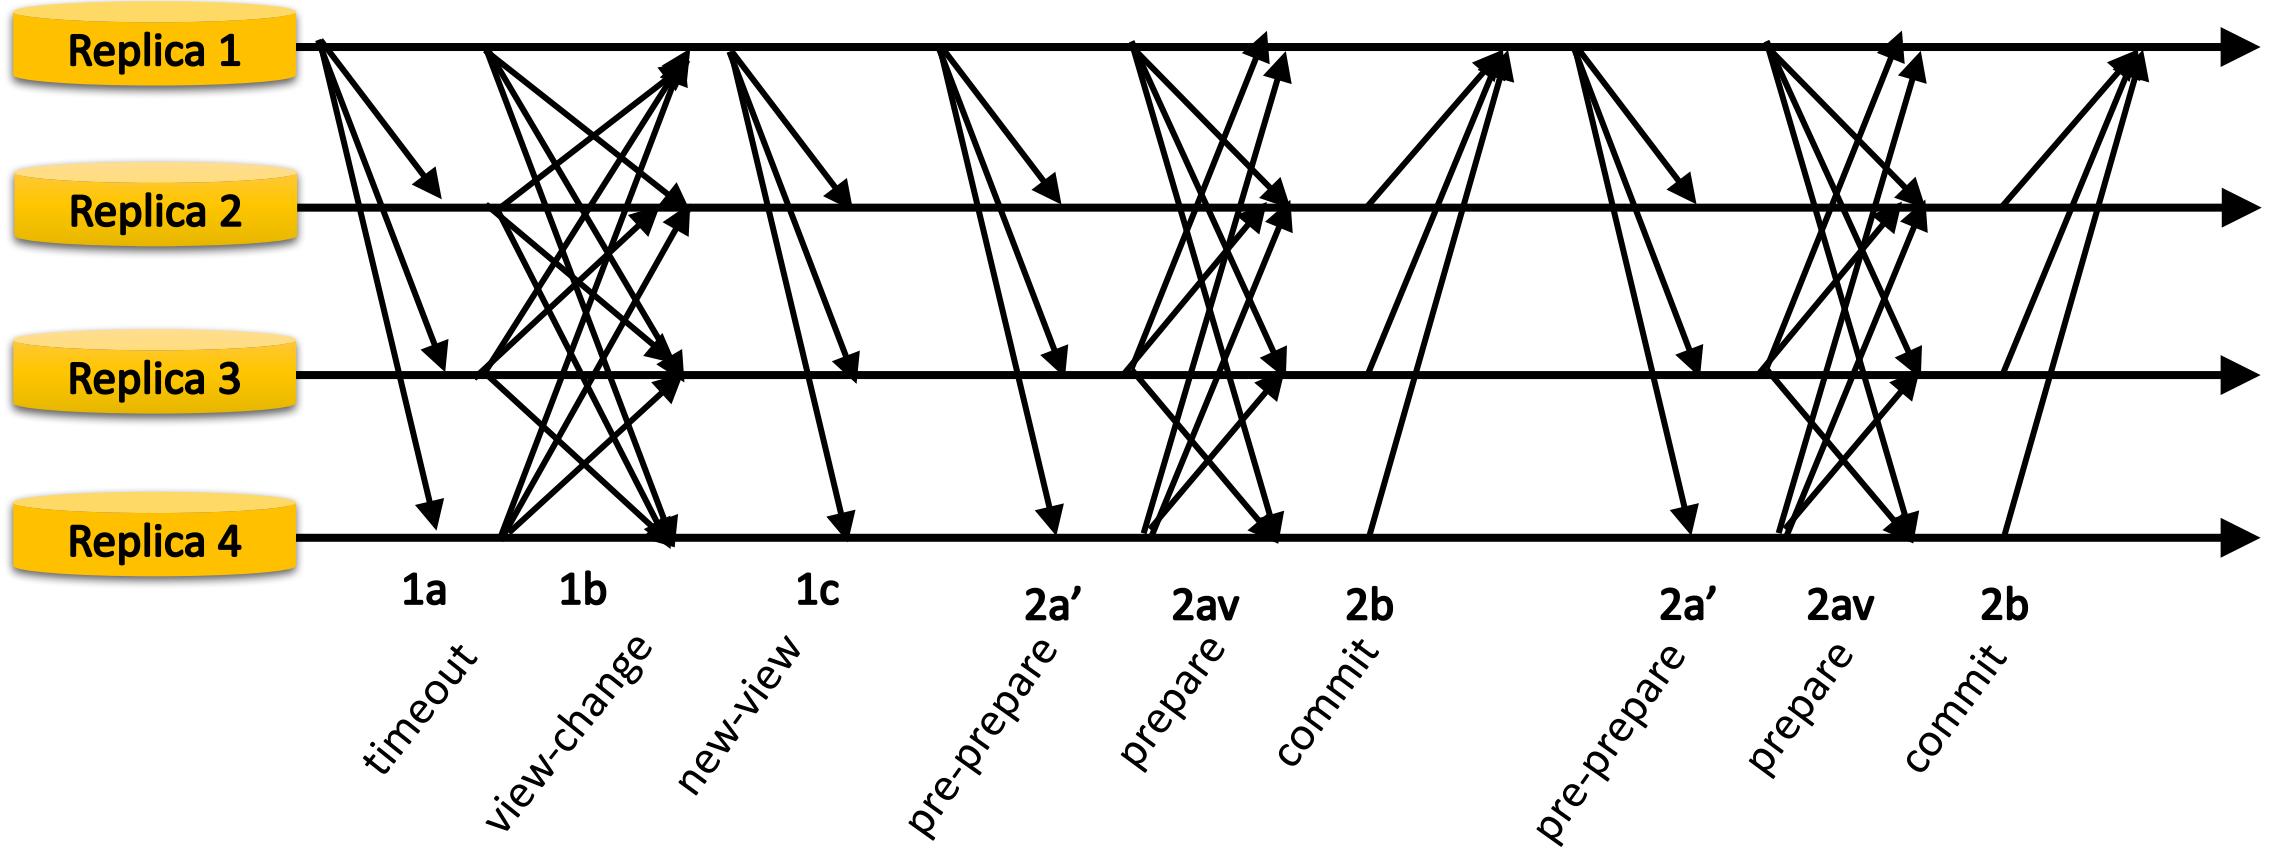


Multi-instance Example, like Multi-Paxos



A little bit renaming, and look what we get here!

Miguel Castro and Barbara Liskov, *Practical Byzantine Fault Tolerance*, OSDI'99



Summary: the refinements

- Paxos + fake acceptors
- Enlarge quorum to tolerate fake non-leaders
- Attach 1b message with previous 2a message to shrink quorum
- Extract 1c from 2a
- Cooperative 2a using 1c
- Broadcast phase 1b to deduct 1c
- Multiple instances with batched phase 1